

COMP9418: Advanced Topics in Statistical Machine Learning

Markov Chains and Hidden Markov Models

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

Instructor: Gustavo Batista

University of New South Wales

Introduction

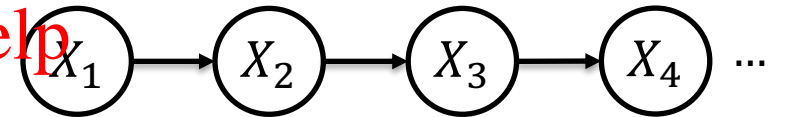
- This lecture discusses two classes of Graphical Models
 - Markov chains
 - Hidden Markov Models (HMM)
- Both models are instances of Dynamic Bayesian Networks (DBN)
 - They have a repeating structure that grows with time or space
 - Such structure is simple and uses the *Markov property*
- The Markov property states that future states are independent of past ones given the current state
 - In Markov chains, all states are observable
 - HMM extend the chains by allowing hidden states
- We will discuss specialised inference algorithms for both classes
 - Applications of these graphical models in domains such as robot localisation

Time and Space

- Several problems require reasoning about sequences

- Such sequences may represent the problem dynamics in time or space

- Examples of applications are speech recognition and robot localization



Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

- Dynamic Bayesian Networks (DBN) allow to incorporate time or space in our models

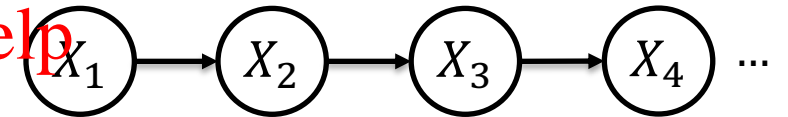
- The two simplest instances of DBNs are Markov chains and Hidden Markov models

Markov Chains

- Markov chain is a *state machine*

- X is a discrete variable and each value is called a *state*
- Transitions between states are nondeterministic

Assignment Project Exam Help



- Parameters

- Prior probabilities $P(X_1)$
- Transition probabilities or *dynamics* $P(X_t | X_{t-1})$

<https://tutorcs.com>

WeChat: cstutorcs

- Stationary assumption

- Transition probabilities are the same for all values of t
- Also known as a *time-homogeneous* chain

Markov Chains: Weather

- States

- $X = \{sun, rain\}$

- Initial distribution

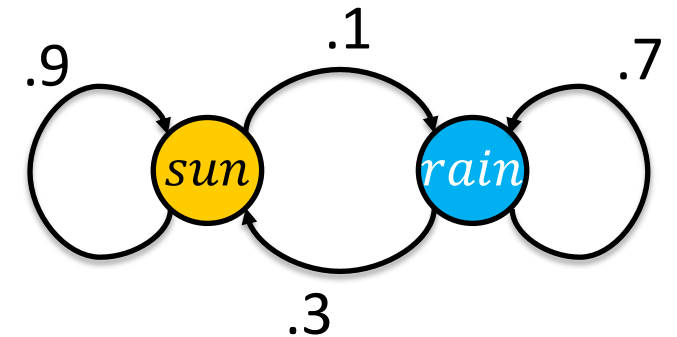
$$X_1 = \begin{pmatrix} 1 & sun \\ 0 & rain \end{pmatrix}$$

- Transition probabilities

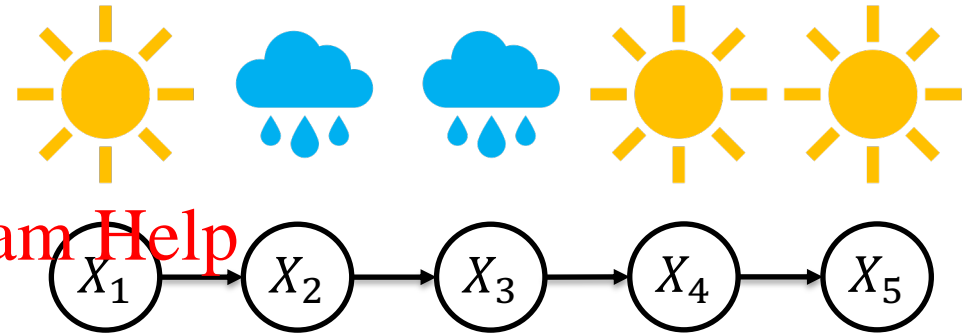
X_{t-1}	X_t	$P(X_t X_{t-1})$
<i>sun</i>	<i>sun</i>	.9
<i>sun</i>	<i>rain</i>	.1
<i>rain</i>	<i>sun</i>	.3
<i>rain</i>	<i>rain</i>	.7

		X_t	
		<i>sun</i>	<i>rain</i>
X_{t-1}	<i>sun</i>	.9	.1
	<i>rain</i>	.3	.7

Matrix of transition probabilities



Transition or state diagram



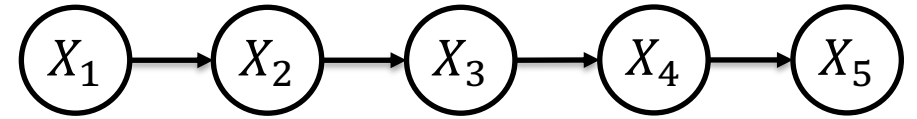
Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

Markov Chains: Independencies

- An relevant question is which independencies are implied in this chain



- We can use d-separation to visually infer the independencies

Assignment Project Exam Help

$$X_1 \perp X_3 | X_2$$

$$X_2 \perp X_4 | X_3$$

- This independence assumption is known as (first order) Markov property

<https://tutorcs.com>

More generally, $X_{t+1} \perp X_{t-1} | X_t$

- Independences are also apparent when we look at the chain rule

WeChat: cstutorcs

- Chain rule if Bayesian networks for this example

$$P(X_1, X_2, X_3, X_4, X_5) = P(X_1)P(X_2|X_1)P(X_3|X_2)P(X_4|X_3)P(X_5|X_4)$$

- Chain rule in general

$$P(X_1, X_2, X_3, X_4, X_5) = P(X_1)P(X_2|X_1)P(X_3|X_2, X_1)P(X_4|X_3, X_2, X_1)P(X_5|X_4, X_3, X_2, X_1)$$

$$X_3 \perp X_1 | X_2$$

$$X_4 \perp X_1, X_2 | X_3$$

$$X_5 \perp X_1, X_2, X_3 | X_4$$

Markov Chains: Independencies

- In general

$$P(X_1, \dots, X_n) = P(X_1) \prod_{i=2}^n P(X_i | X_{i-1})$$

Assignment Project Exam Help

$|X|^n$ parameters

$|X| + (n - 1)|X|^2$ parameters

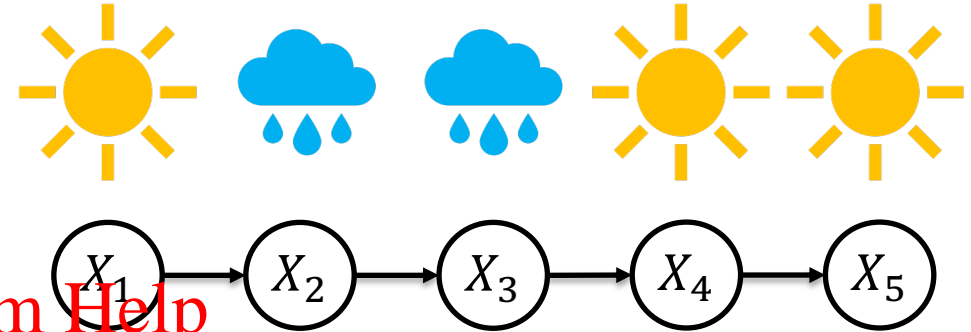
<https://tutorcs.com>

WeChat: cstutorcs

$|X| + |X|^2$ parameters

$$X_1 = \begin{pmatrix} 1 & \text{sun} \\ 0 & \text{rain} \end{pmatrix}$$

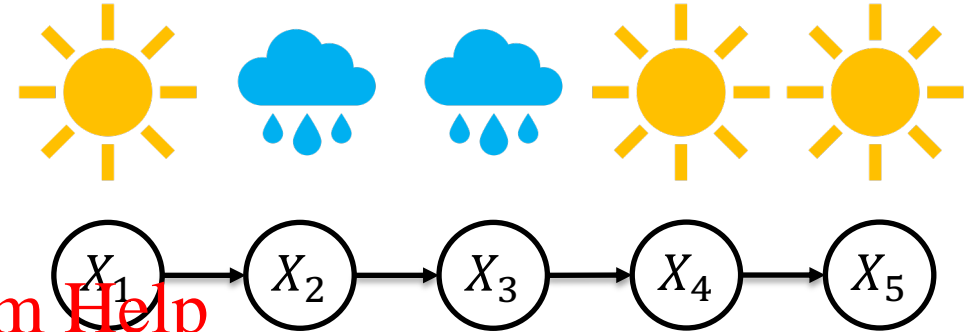
X_{t-1}	X_t	$P(X_t X_{t-1})$
sun	sun	.9
sun	rain	.1
rain	sun	.3
rain	rain	.7



We also assume that $P(X_t | X_{t-1})$ is the same for all t (stationarity)

Probability of a State Sequence

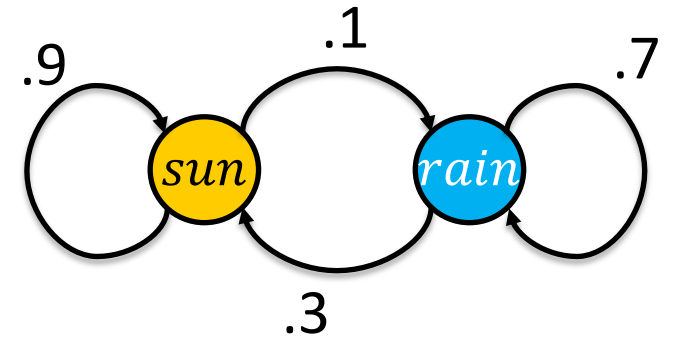
- The probability of a sequence is the product of the transition probabilities
 - This comes directly from the chain rule



$$P(X_1, X_2, X_3, X_4, X_5) = P(X_1)P(X_2|X_1)P(X_3|X_2)P(X_4|X_3)P(X_5|X_4)$$

For example, what is the probability of the sequence: sun, rain, rain, sun, sun?

$$P(\text{sun}, \text{rain}, \text{rain}, \text{sun}, \text{sun}) = 1(.1)(.7)(.3)(.9) = .189$$



$$X_1 = \begin{pmatrix} 1 & \text{sun} \\ 0 & \text{rain} \end{pmatrix}$$

Probability of Staying in a Certain State

- The probability of staying in a certain state for d steps

- It is the probability of a sequence in this state for $d - 1$ steps then going to a different state

Assignment Project Exam Help

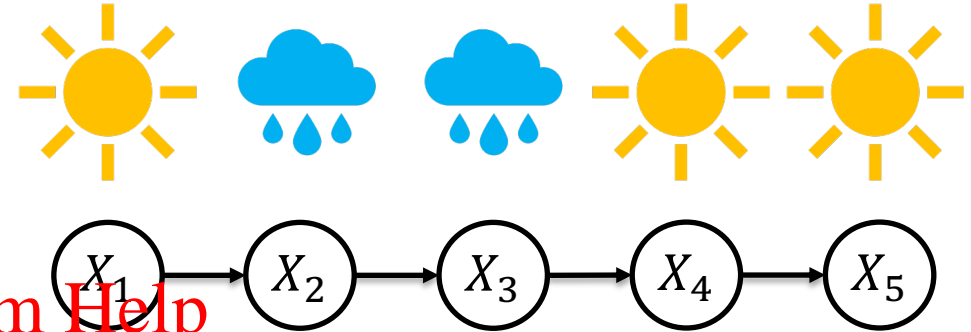
<https://tutorcs.com>

WeChat: cstutorcs

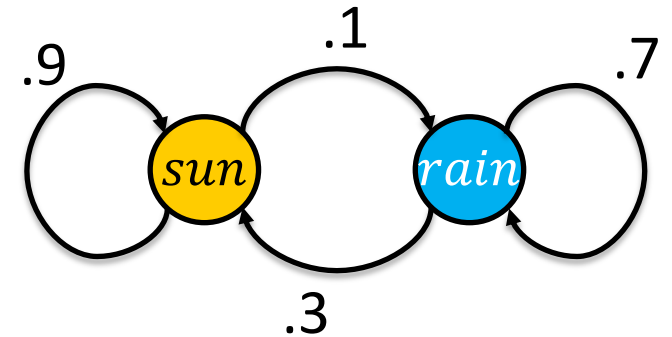
$$\mathbf{S}_s^d = \{X_i = s : 1 \leq i \leq d\}$$

$$P(\mathbf{S}_s^d) = P(s|s)^{d-1}(1 - P(s|s))$$

$$P(\mathbf{S}_{rain}^3) = P(rain|rain)^{3-1}(1 - P(rain|rain)) = (.7^2)(1 - .7) = .147$$



For example, what is the probability of three raining days?



$$X_1 = \begin{pmatrix} 1 & \text{sun} \\ 0 & \text{rain} \end{pmatrix}$$

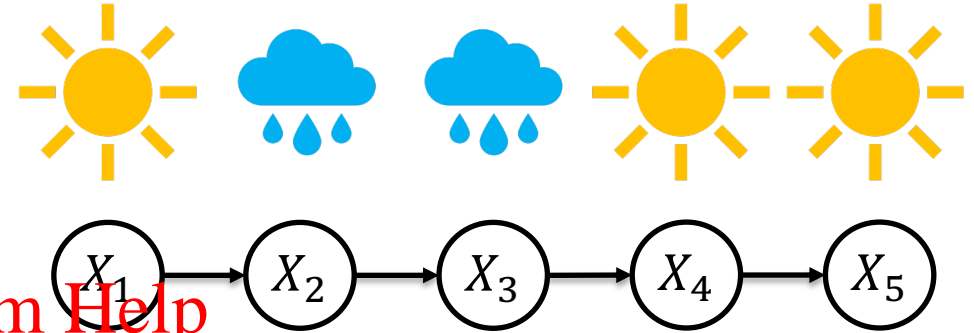
Expected Time in a State

- The average duration of a sequence is a certain state
 - It is the expected number of time steps in that state

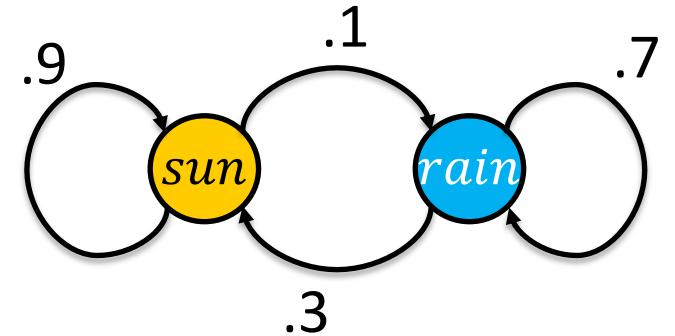
$$\begin{aligned}\mathbb{E}[S_s] &= \sum_i^{\infty} P(s_s^i) i \\ &= \sum_i^{\infty} i P(s|s)^{i-1} (1 - P(s|s)) \\ &= \frac{1}{1 - P(s|s)} \\ \mathbb{E}(S_{rain}) &= \frac{1}{1 - .7} = 3.33\end{aligned}$$

Assignment Project Exam Help

<https://tutorcs.com>
WeChat: cstutorcs



For example, what is the expected number of raining days?



$$X_1 = \begin{pmatrix} 1 & sun \\ 0 & rain \end{pmatrix}$$

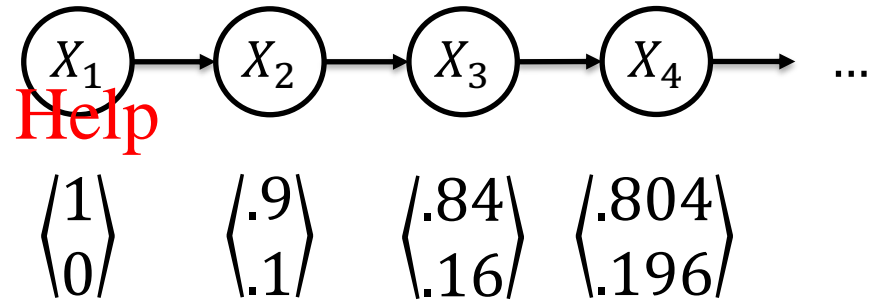
Mini-Forward Algorithm

- What is $P(X)$ on some day t ?
 - We can obtain an answer by simulating the chain

Assignment Project Exam Help

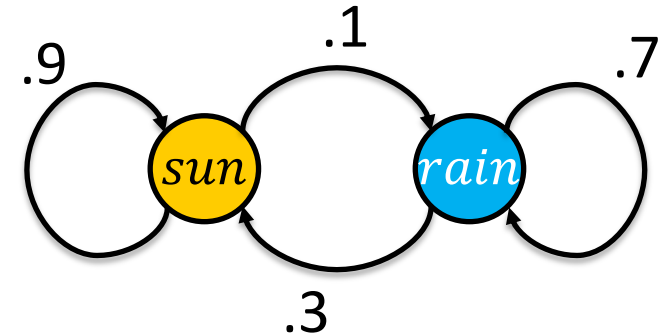
<https://tutorcs.com>

WeChat: cstutorcs



$P(x_1)$ is known

$$P(x_t) = \sum_{x_{t-1}} P(x_t, x_{t-1})$$
$$= \sum_{x_{t-1}} P(x_t | x_{t-1}) P(x_{t-1})$$



Mini-Forward Algorithm

Input: time n , transition probability $P(X_t|X_{t-1})$, prior probability of states $P(X_1)$

Output: $P(X_n)$

for each state x **do**

$p[x, 1] \leftarrow P(X_1 = x)$

for $t \leftarrow 2$ **to** n **do**

for each state x_t **do**

$p[x_t, t] = 0$

for each state x_{t-1} **do**

$p[x_t, t] \leftarrow p[x_t, t] + p[x_{t-1}, t-1]P(x_t|x_{t-1})$

return $p[x, n]$

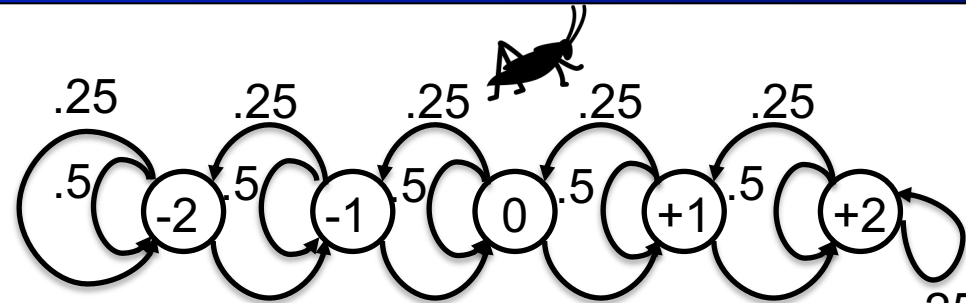
Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

$$O(n|X|^2)$$

Grasshopper Example



Assignment Project Exam Help

	-2	-1	0	1	2
$P(X_1)$	0	0	1	0	0
$P(X_2)$	0	.25	.5	.25	0
$P(X_3)$	$.25^2 = .0625$	$2(.5)(.25) = .25$	$.5^2 + 2(.25)^2 = .375$	$2(.5)(.25) = .25$	$.25^2 = .0625$

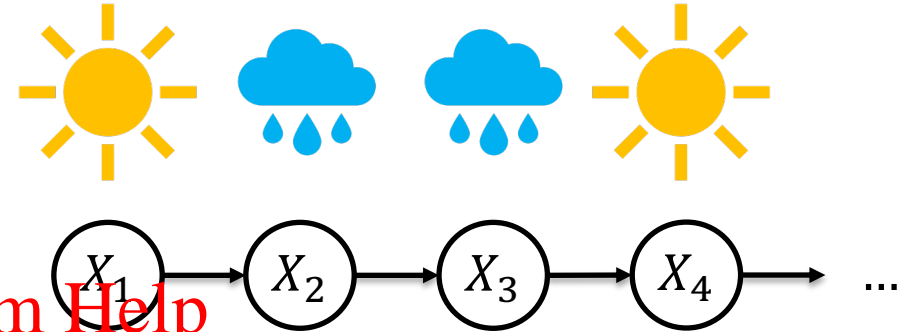
$$P(X_t) = P(X_{t-1})T \quad [0 \ 0 \ 1 \ 0 \ 0] \begin{bmatrix} .75 & .25 & & & \\ .25 & .5 & .25 & & \\ & .25 & .5 & .25 & \\ & & .25 & .5 & .25 \\ & & & .25 & .75 \end{bmatrix} = [0 \ .25 \ .5 \ .25 \ 0]$$

Stationary Distributions

- Starting with a sunny day

$$\begin{pmatrix} 1 \\ 0 \end{pmatrix} \quad \begin{pmatrix} .9 \\ .1 \end{pmatrix} \quad \begin{pmatrix} .84 \\ .16 \end{pmatrix} \quad \begin{pmatrix} .804 \\ .196 \end{pmatrix} \quad \dots \quad \begin{pmatrix} .75 \\ .25 \end{pmatrix}$$

Assignment Project Exam Help



- Starting with a rainy day

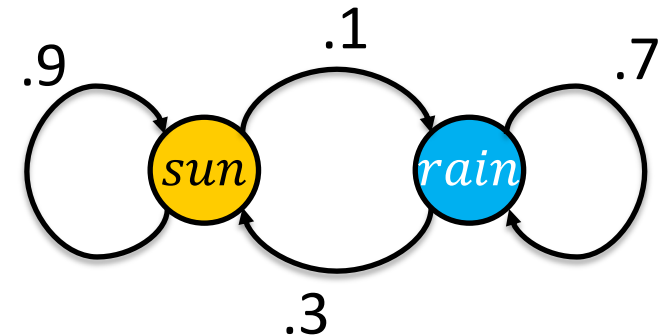
$$\begin{pmatrix} 0 \\ 1 \end{pmatrix} \quad \begin{pmatrix} .3 \\ .7 \end{pmatrix} \quad \begin{pmatrix} .48 \\ .52 \end{pmatrix} \quad \begin{pmatrix} .588 \\ .412 \end{pmatrix} \quad \dots \quad \begin{pmatrix} .75 \\ .25 \end{pmatrix}$$

<https://tutorcs.com>

WeChat: cstutorcs

- Starting with an unknown day

$$\begin{pmatrix} p \\ 1 - p \end{pmatrix} \quad \dots \quad \begin{pmatrix} .75 \\ .25 \end{pmatrix}$$



Stationary Distributions

- For most chains

- Influence of the initial distribution gets less and less over time

- The distribution we end up in is independent of the initial distribution

Assignment Project Exam Help

<https://tutorcs.com>

- Stationary distribution

WeChat: cstutorcs

- The *stationary distribution* π of the chain is the distribution we obtain if the chain converges
- The stationary distribution satisfies

$$P_{\infty}(X) = P_{\infty+1}(X) = \sum_x P(X|x)P_{\infty}(x)$$

$$\pi(X) = \sum_x P(X|x)\pi(x)$$

$$\pi = \pi T$$

Stationary Distributions

- Question: What is $P(X_\infty)$?

$$\pi(\text{sun}) = P(\text{sun}|\text{sun})\pi(\text{sun}) + P(\text{sun}|\text{rain})\pi(\text{rain})$$

$$\pi(\text{rain}) = P(\text{rain}|\text{sun})\pi(\text{sun}) + P(\text{rain}|\text{rain})\pi(\text{rain})$$

$$\pi(\text{sun}) = 0.9 \pi(\text{sun}) + 0.3 \pi(\text{rain})$$

$$\pi(\text{rain}) = 0.1 \pi(\text{sun}) + 0.7 \pi(\text{rain})$$

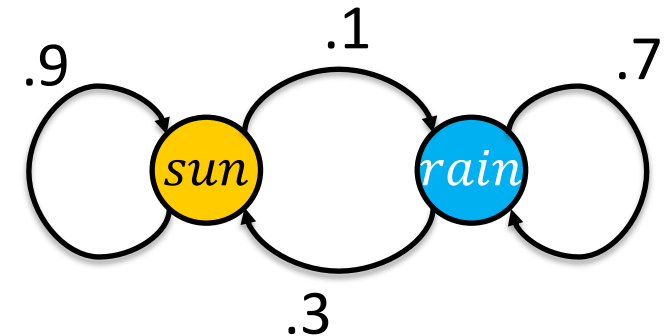
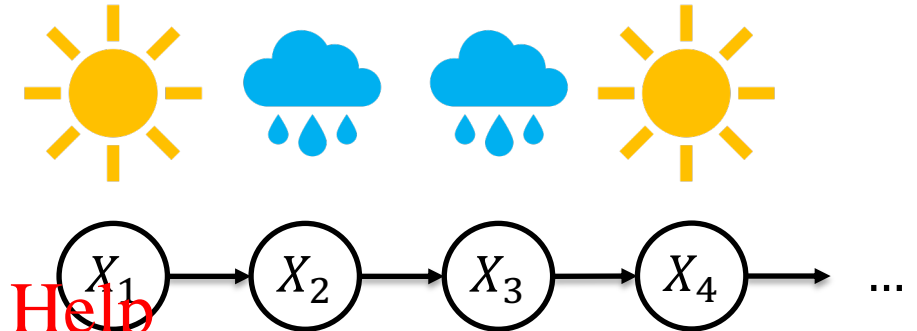
$$\pi(\text{sun}) = 3\pi(\text{rain})$$

$$\pi(\text{rain}) = 1/3\pi(\text{sun})$$

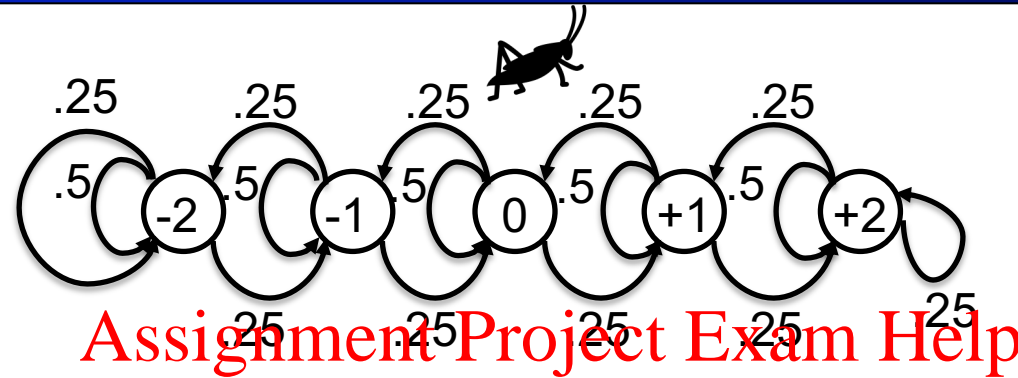
$$\pi(\text{sun}) + \pi(\text{rain}) = 1$$

$$\pi(\text{sun}) = 3/4$$

$$\pi(\text{rain}) = 1/4$$



Stationary Distributions: Grasshopper



<https://tutorcs.com>

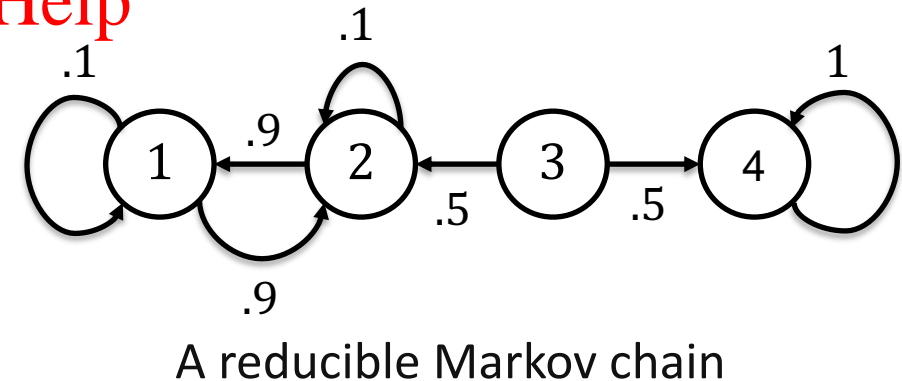
WeChat: cstutorcs

- What is the stationary distribution?

$$T = \begin{bmatrix} .75 & .25 & & & \\ .25 & .5 & .25 & & \\ & .25 & .5 & .25 & \\ & & .25 & .5 & .25 \\ & & & .25 & .75 \end{bmatrix}$$

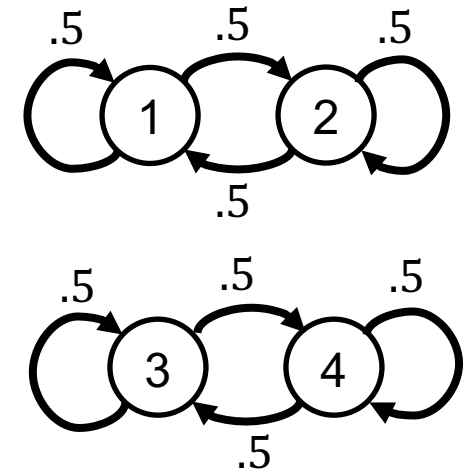
Irreducible Markov Chains

- A Markov chain is *irreducible* if every state x' is reachable from every other state x
 - That is, for every pair of states x and x' , there is some time t such that the $P(X_t = x' | X_1 = x) > 0$
 - Also known as *regular* or *ergodic* chain
- In this case, the states of the Markov chain are said to be *recurrent*
 - Each state is guaranteed to be visited an infinite number of times when we simulate the chain



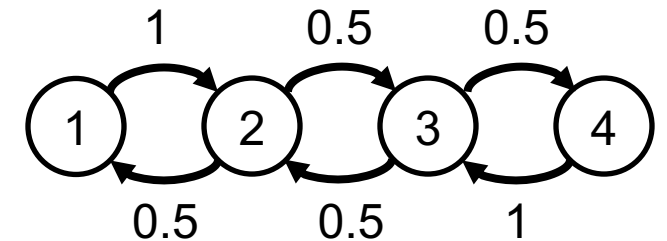
Stationary Distribution

- Every (finite state) Markov chain has at least one stationary distribution
 - Yet an irreducible Markov chain is guaranteed to have a unique stationary distribution
- An irreducible chain may or may not converge to its stationary distribution
 - To guarantee convergence, we need an additional property: *Aperiodicity*



Aperiodic Markov Chains

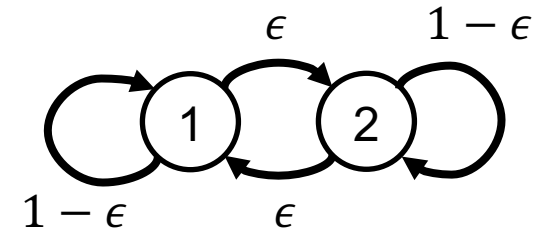
- A Markov chain is aperiodic if it is possible to return to any state at any time
 - There exists an t such that for all state x and all $t' \geq t$,
 $P(X_{t'} = x \mid X_1 = x) > 0$
- An irreducible and aperiodic Markov chain converges to a unique stationary distribution
 - Irreducible: we can go from any state to any state
 - Aperiodic: avoids chains that alternates forever between states without ever settling in a stationary distribution



An irreducible but periodic Markov chain

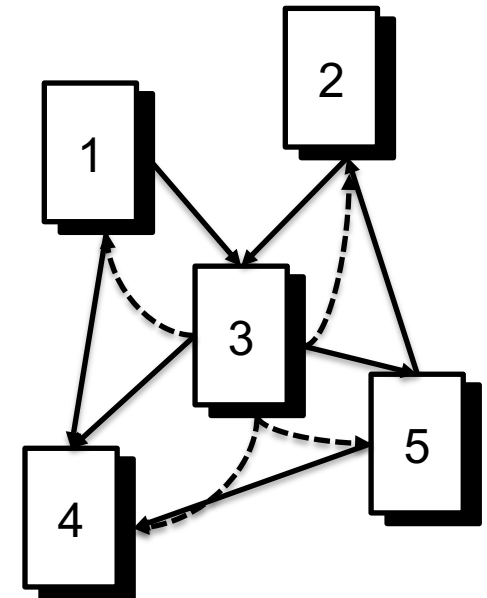
Markov Chains Convergence

- Although an irreducible and aperiodic Markov chain converges to a single stationary distribution, the convergence can be slow
- In this example, the stationary distribution is close to $(0.5, 0.5)$
- For a small ϵ it will take a very long time to reach the stationary distribution
- We stay in the same state with high probability and rarely transition to another state
- The average of these states will converge to $(0.5, 0.5)$, but the convergence will be very slow



Applications of Markov Chains

- Markov chains have several well-known applications
 - Markov chain Monte Carlo (MCMC) is a powerful approximate inference algorithm used in statistical software such as Stan
 - MCs are part of the (LZMA) Lempel-Ziv-Markov compression algorithm used in 7zip
 - PageRank algorithm used by Google 1.0 is a direct application of MCs
- PageRank
 - Model the web as a state graph: pages are states and hyperlinks are transitions
 - Each transition from state i has a probability $\frac{\alpha}{k_i}$, where α is a constant parameter and k_i is the outgoing degree of node i
 - Compute a stationary distribution. But it is not unique. Why?
 - Augment the graph with phantom transitions of weight $\frac{1-\alpha}{N}$, where N is the number of nodes in the graph

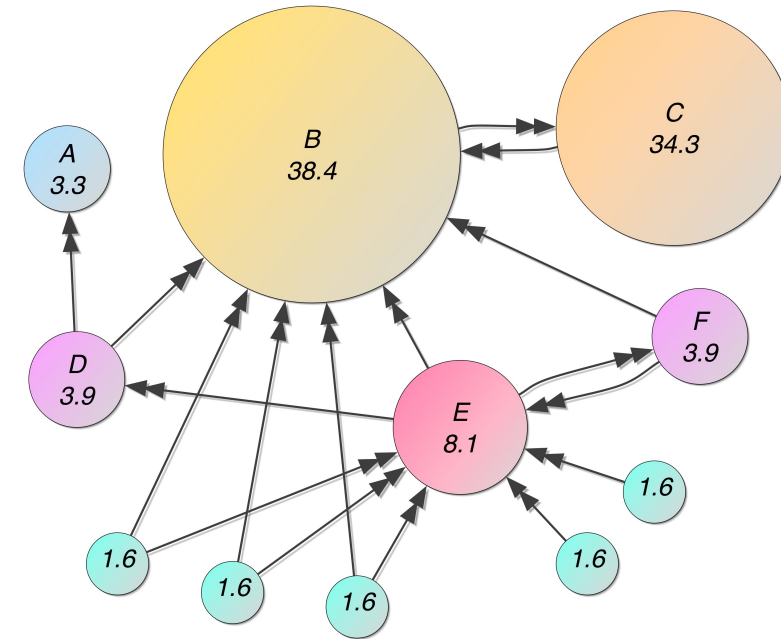


Applications of Markov Chains

- Markov chains have several well-known applications
 - Markov chain Monte Carlo (MCMC) is a powerful approximate inference algorithm used in statistical software such as Stan
 - MCs are part of the (LZMA) Lempel-Ziv-Markov compression algorithm used in 7zip
 - PageRank algorithm used by Google 1.0 is a direct application of MCs

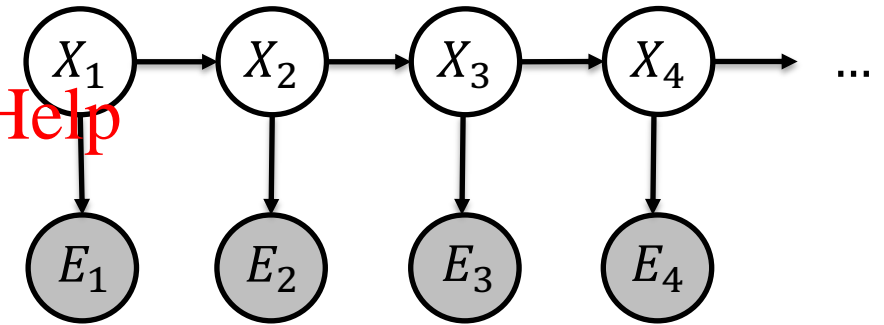
- PageRank

- Model the web as a state graph: pages are states and hyperlinks are transitions
- Each transition from state i has a probability $\frac{\alpha}{k_i}$, where α is a constant parameter and k_i is the outgoing degree of node i
- Compute a stationary distribution. But it is not unique. Why?
- Augment the graph with phantom transitions of weight $\frac{1-\alpha}{N}$, where N is the number of nodes in the graph



Hidden Markov Models (HMM)

- Hidden Markov Models (HMM) are Markov chains where the states are not directly observable
 - In the weather example, the weather may not be directly observable
 - Instead, we use sensors, such as temperature, air pressure, humidity, wind speed, etc.
- HMM has two components
 - Underlying Markov chain over states X
 - Observable outputs (effects of the states) at each time step
 - These outputs are often called *emissions*



Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

HMM Weather Example

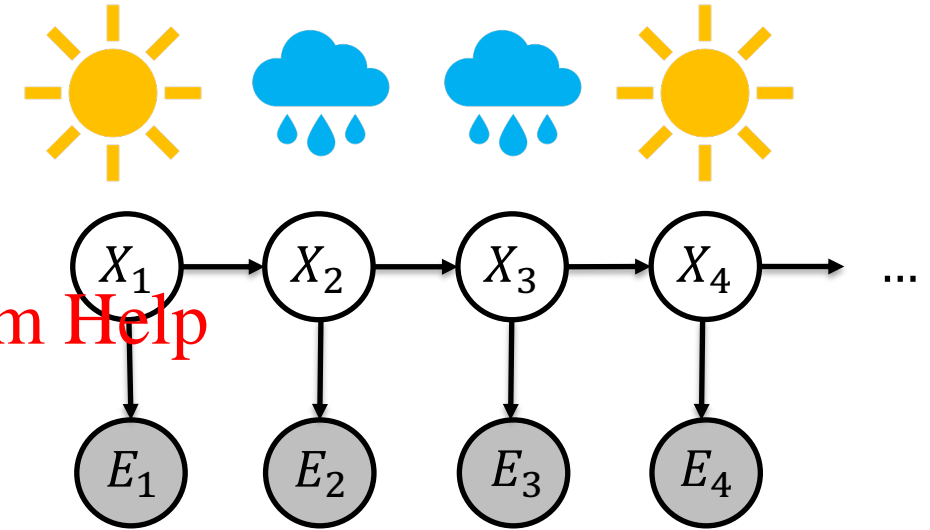
- HMM parameters

- Initial distribution $P(X_1)$
- Transition probabilities $P(X_t|X_{t-1})$
- Emission probabilities $P(E_t|X_t)$

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs



X_1	$P(X_1)$
<i>sun</i>	.5
<i>rain</i>	.5

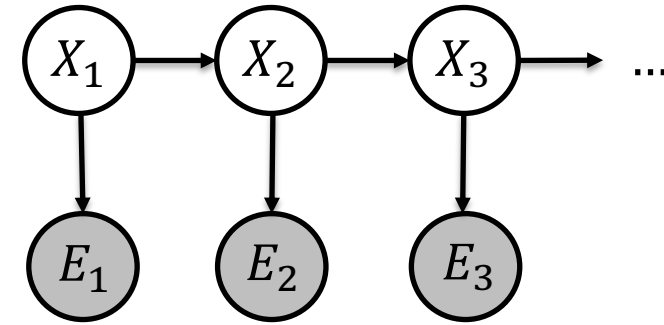
X_{t-1}	X_t	$P(X_t X_{t-1})$
<i>sun</i>	<i>sun</i>	.7
<i>sun</i>	<i>rain</i>	.3
<i>rain</i>	<i>sun</i>	.3
<i>rain</i>	<i>rain</i>	.7

X_t	E_t	$P(E_t X_t)$
<i>sun</i>	<i>umb</i>	.2
<i>sun</i>	\overline{umb}	.8
<i>rain</i>	<i>umb</i>	.9
<i>rain</i>	\overline{umb}	.1

HMM: Independencies

- The chain rule of Bayesian networks for HMMs

$$P(X_1, E_1, \dots, X_n, E_n) = P(X_1)P(E_1|X_1) \prod_{t=1}^n P(X_t|X_{t-1})P(E_t|X_t)$$



- Independences are also apparent when we look at the chain rule

<https://tutorcs.com>

- Chain rule for Bayesian networks for this example

$$P(X_1, E_1, X_2, E_2, X_3, E_3) = P(X_1)P(E_1|X_1)P(X_2|X_1)P(E_2|X_2)P(X_3|X_2)P(E_3|X_3)$$

- Chain rule in general

$$P(X_1, E_1, X_2, E_2, X_3, E_3) = P(X_1)P(E_1|X_1)P(X_2|X_1, E_1)P(E_2|X_2, X_1, E_1)P(X_3|X_2, X_1, E_2, E_1)P(E_3|X_3, X_2, X_1, E_2, E_1)$$

$$X_2 \perp E_1 | X_1$$

$$X_3 \perp X_1, E_1, E_2 | X_2$$

$$E_2 \perp X_1, E_1 | X_2$$

$$E_3 \perp X_1, X_2, E_1, E_2 | X_3$$

HMM: Independencies

- In general, HMM have the following independency assumptions

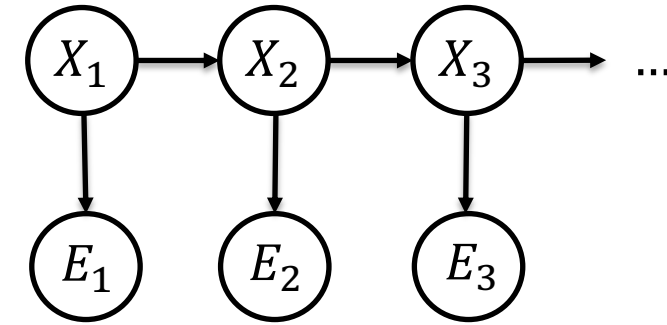
- A state is independent of all past states and all past evidence given the previous state (Markov property)

$$X_t \perp X_1, \dots, X_{t-2}, E_1, \dots, E_{t-1} | X_{t-1}$$

- Evidence is independent of all past evidence and all past states given the current state (independence of observations)

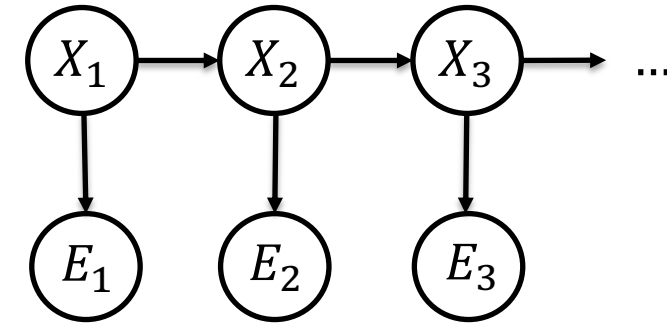
$$E_t \perp X_1, \dots, X_{t-1}, E_1, \dots, E_{t-1} | X_t$$

- Transition and emission probabilities are the same for all values of t (stationary process)



HMM: Inference

- We start with a first task of tracking the distribution $P(X_t)$ over time
 - This task is known as *filtering or monitoring*
 - We use $B(X_t) = P(X_t|e_1, \dots, e_t)$ to denote the *belief of state*
 - We start with $B(X_1)$, usually using a uniform distribution
 - Update $B(X_t)$ as time passes and we get new observations
- The inference has two main steps
 - Passage of time
 - Observation



Passage of Time

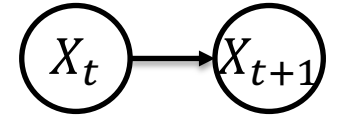
- Suppose we know the current state of belief $B(X_t)$

$$B(X_t) = P(X_t|e_{1:t})$$

- We need to update it as one unit of time passes. Our aim is to compute $P(X_{t+1}|e_{1:t})$

<https://tutorcs.com>

$$\begin{aligned} P(X_{t+1}|e_{1:t}) &= \sum_{x_t} P(X_{t+1}, x_t|e_{1:t}) \\ &= \sum_{x_t} P(X_{t+1}|x_t, e_{1:t})P(x_t|e_{1:t}) \\ &= \sum_{x_t} P(X_{t+1}|x_t)P(x_t|e_{1:t}) \\ &= \sum_{x_t} P(X_{t+1}|x_t)B(x_t) \end{aligned}$$



Observation

- Given we updated the belief with passage of time

- We know $P(X_{t+1}|e_{1:t})$ and we need to update it to

$$B(X_{t+1}) = P(X_{t+1}|e_{1:t+1})$$

Assignment Project Exam Help

$$P(X_{t+1}|e_{1:t+1}) = P(X_{t+1}|e_{t+1}, e_{1:t})$$

$$= P(X_{t+1}, e_{t+1}|e_{1:t}) / P(e_{t+1}|e_{1:t})$$

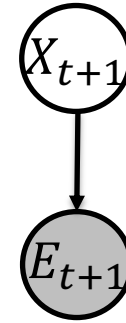
$$\propto P(X_{t+1}, e_{t+1}|e_{1:t})$$

$$= P(e_{t+1}, X_{t+1}|e_{1:t})$$

$$= P(e_{t+1}|X_{t+1}, e_{1:t}) P(X_{t+1}|e_{1:t})$$

$$= P(e_{t+1}|X_{t+1}) P(X_{t+1}|e_{1:t})$$

$$B(X_{t+1}) \propto P(e_{t+1}|X_{t+1}) P(X_{t+1}|e_{1:t})$$



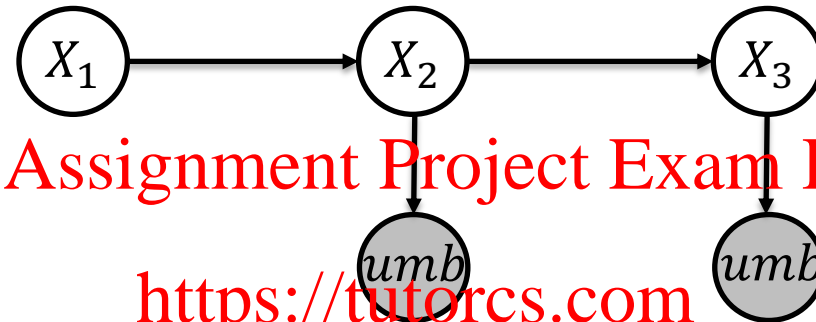
We must renormalise
the results by $\sum B(X_{t+1})$

HMM Weather Example

$$\begin{aligned} B(\text{sun}) &= .5 \\ B(\text{rain}) &= .5 \end{aligned}$$

$$\begin{aligned} B(\text{sun}) &= .5 \\ B(\text{rain}) &= .5 \end{aligned}$$

$$\begin{aligned} B(\text{sun}) &= .373 \\ B(\text{rain}) &= .627 \end{aligned}$$



Assignment Project Exam Help

<https://tutorcs.com>

WeChat: estutorcs

$$B(\text{sun}) = .182$$

$$B(\text{sun}) = .117$$

$$B(\text{rain}) = .818$$

$$B(\text{rain}) = .883$$

X_1	$P(X_1)$
<i>sun</i>	.5
<i>rain</i>	.5

X_{t-1}	X_t	$P(X_t X_{t-1})$
<i>sun</i>	<i>sun</i>	.7
<i>sun</i>	<i>rain</i>	.3
<i>rain</i>	<i>sun</i>	.3
<i>rain</i>	<i>rain</i>	.7

X_t	E_t	$P(E_t X_t)$
<i>sun</i>	<i>umb</i>	.2
<i>sun</i>	$\overline{\text{umb}}$.8
<i>rain</i>	<i>umb</i>	.9
<i>rain</i>	$\overline{\text{umb}}$.1

Forward Algorithm

- Suppose we have a sequence of evidence observations and we want to know the state belief at the end of the sequence

$$\begin{aligned}
 B(X_t) &= P(X_t | e_{1:t}) \\
 P(X_t | e_{1:t}) &\propto P(X_t, e_{1:t}) \\
 &= \sum_{x_{t-1}} P(X_t, x_{t-1}, e_{1:t}) \\
 &= \sum_{x_{t-1}} P(X_t, x_{t-1}, e_t, e_{1:t-1}) \\
 &= \sum_{x_{t-1}} P(x_{t-1}) P(X_t | x_{t-1}) P(e_t | x_{t-1}, X_t) P(e_{1:t-1} | e_t, x_{t-1}, X_t) \\
 &= \sum_{x_{t-1}} P(x_{t-1}) P(X_t | X_{t-1}) P(e_t | X_t) P(e_{1:t-1} | x_{t-1}) \\
 &= \sum_{x_{t-1}} P(X_t | x_{t-1}) P(e_t | X_t) P(e_{1:t-1}, x_{t-1}) \\
 &= P(e_t | X_t) \sum_{x_{t-1}} P(X_t | x_{t-1}) P(x_{t-1}, e_{1:t-1})
 \end{aligned}$$

You can renormalise every step, but this algorithm often renormalised only the final one

Forward Algorithm

Input: time n , transition probability T , emission probability E , prior probability of states $P(X_1)$, sequence of observations $\{e_2, \dots, e_t\}$

Output: $B(X_t)$

for each state x **do**

$p[x, 1] \leftarrow P(X_1 = x)$

for $t \leftarrow 2$ to n **do**

for each state x_t **do**

$p[x_t, t] = 0$

for each state x_{t-1} **do**

$p[x_t, t] \leftarrow p[x_t, t] + p[x_{t-1}, t-1]T(x_t|x_{t-1})$

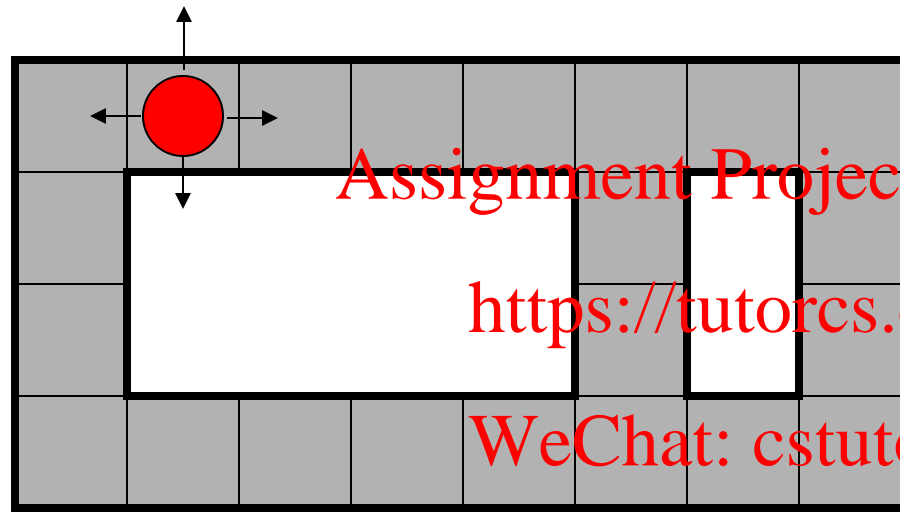
$p[x_t, t] \leftarrow p[x_t, t]E(e_t|x_t)$

return normalised $p[x, n]$ for all states x

$$O(n|X|^2)$$

Example: Robot Localization

Example from
Michael Pfeiffer



Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

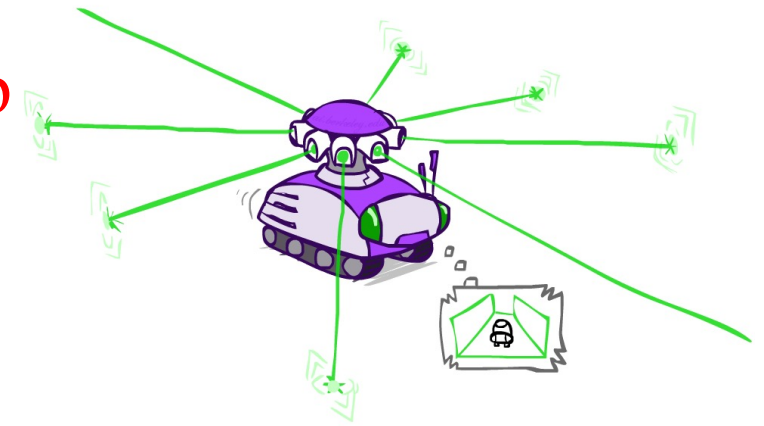


Prob

0

1

$t=0$

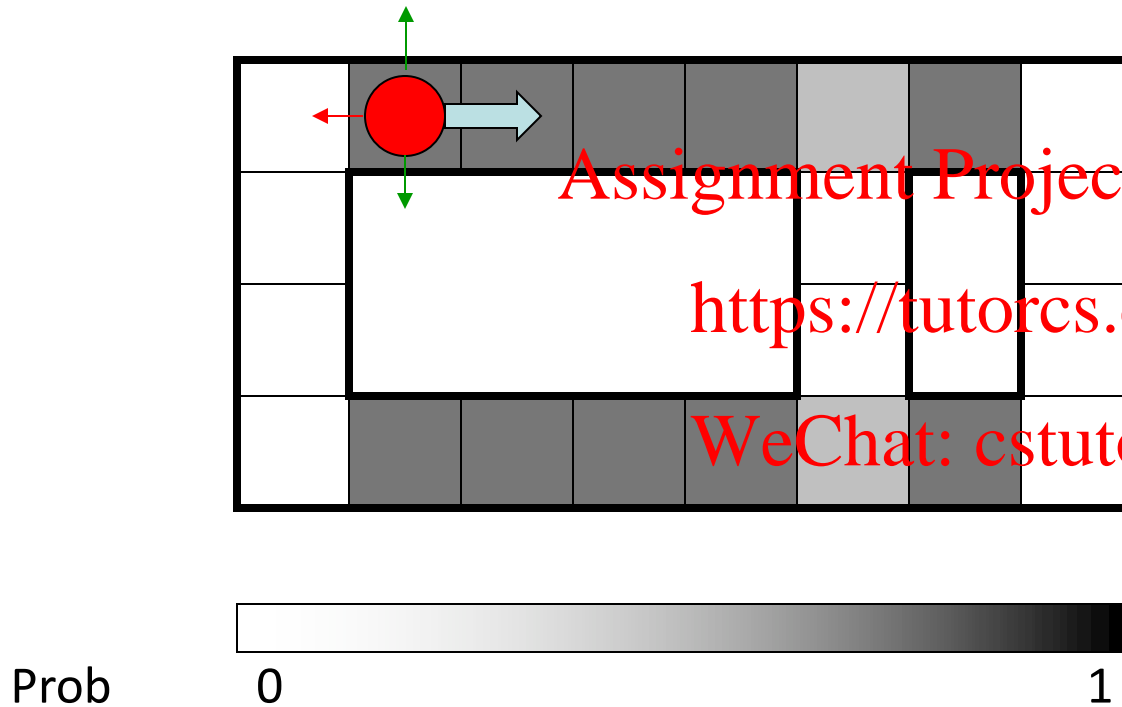


Sensor model: can read in which directions there is a wall,
never more than 1 mistake

Motion model: may not execute action with small prob.

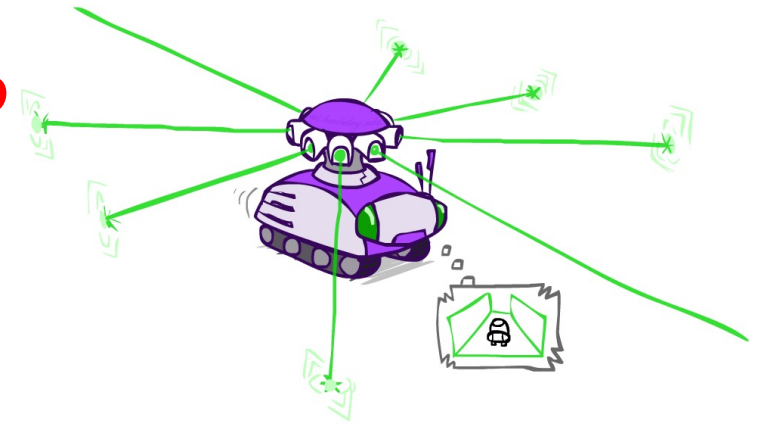
Slide from Berkeley AI course

Example: Robot Localization



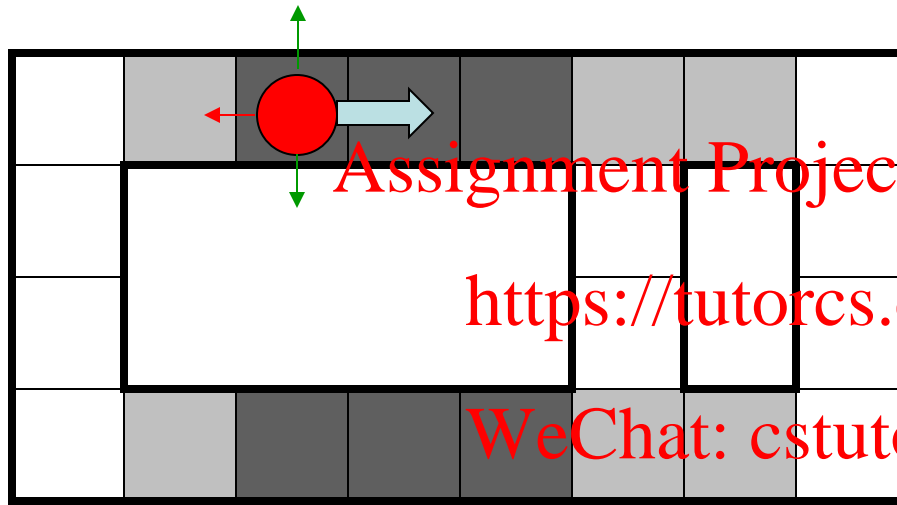
$t=1$

Lighter grey: was possible to get the reading, but less likely b/c required 1 mistake



Slide from Berkeley AI course

Example: Robot Localization

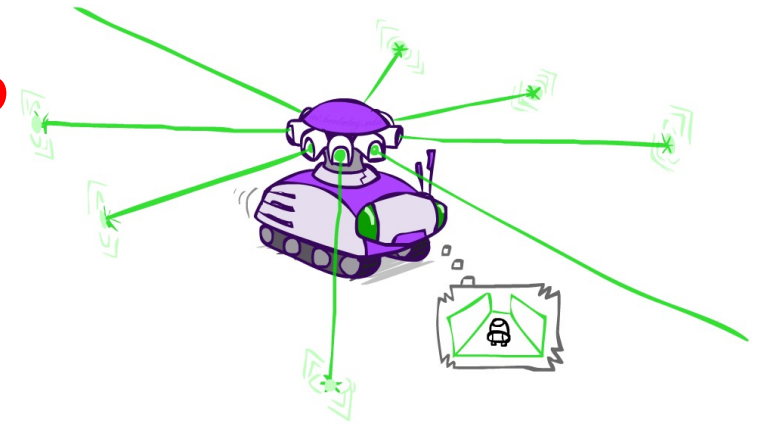


Prob

0

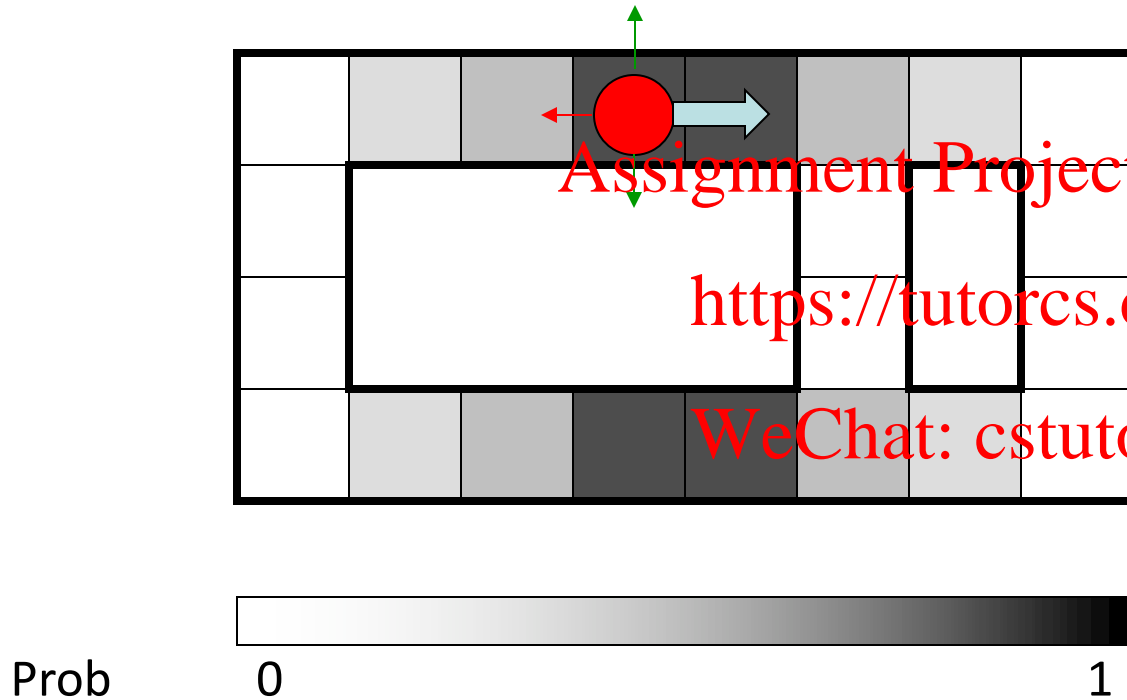
1

$t=2$



Slide from Berkeley AI course

Example: Robot Localization

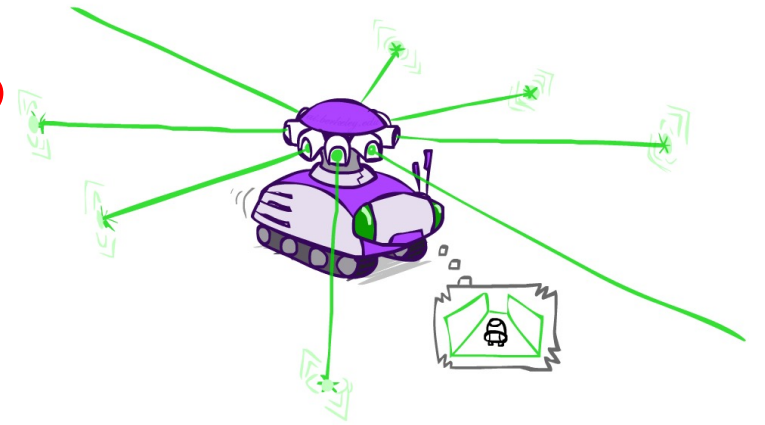


Assignment Project Exam Help

<https://tutorcs.com>

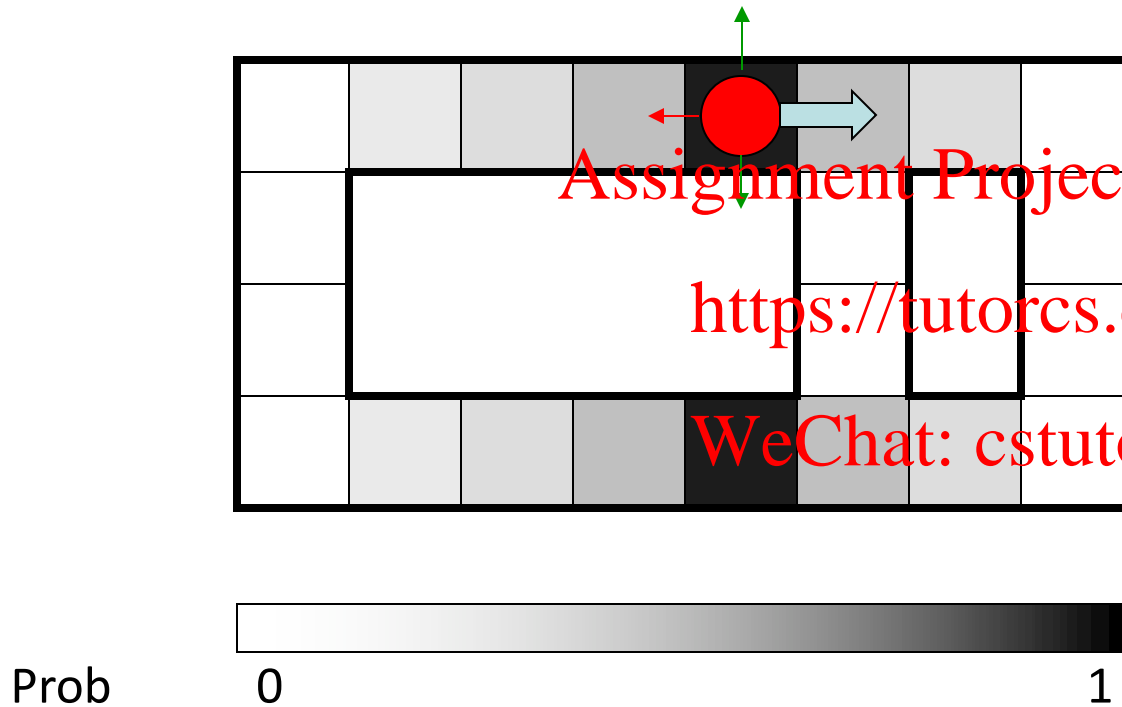
WeChat: cstutorcs

$t=3$



Slide from Berkeley AI course

Example: Robot Localization

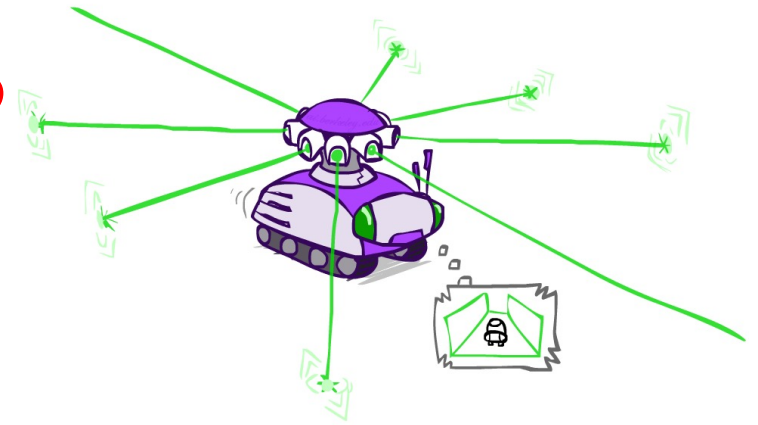


Assignment Project Exam Help

<https://tutorcs.com>

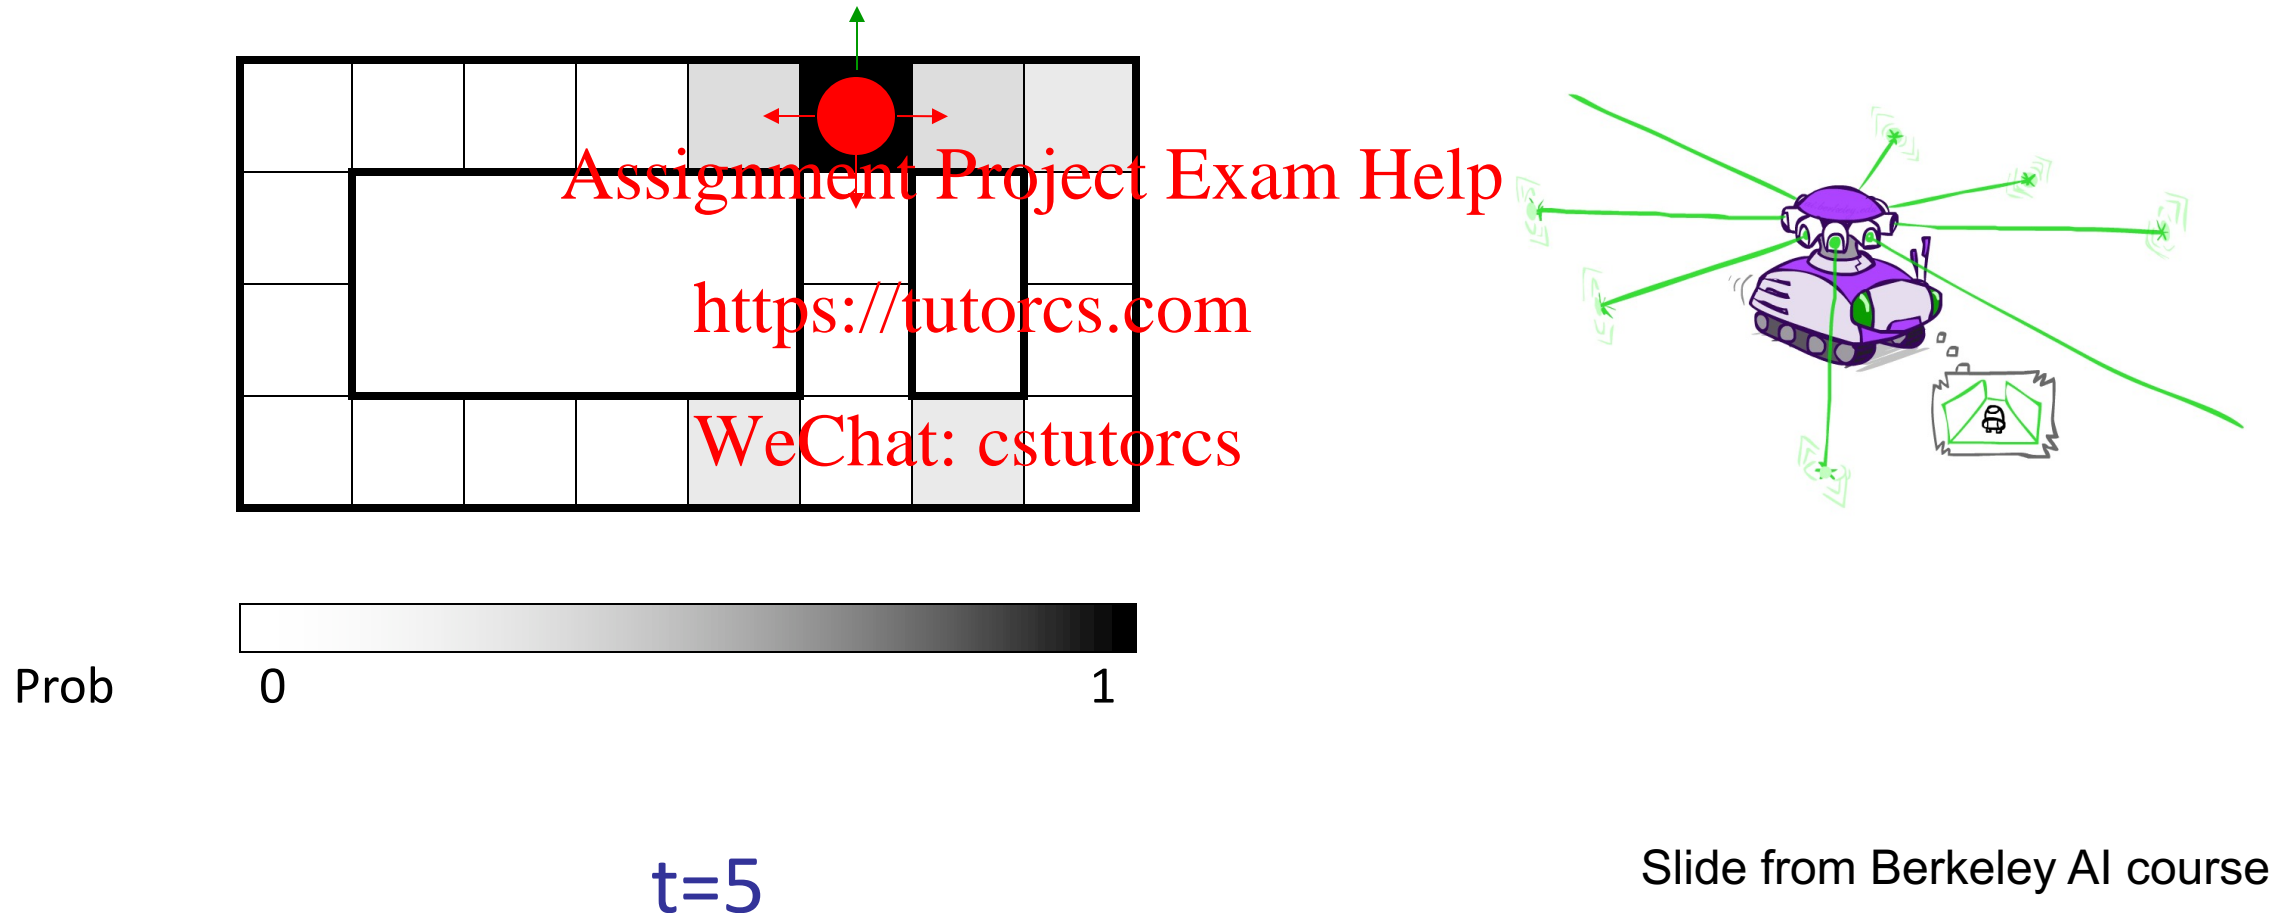
WeChat: cstutorcs

$t=4$



Slide from Berkeley AI course

Example: Robot Localization



Most Probable Explanation (MPE)

- The forward algorithm tracks the probability of the states

- These probabilities are updated with as time passes and we observe evidence

Assignment Project Exam Help

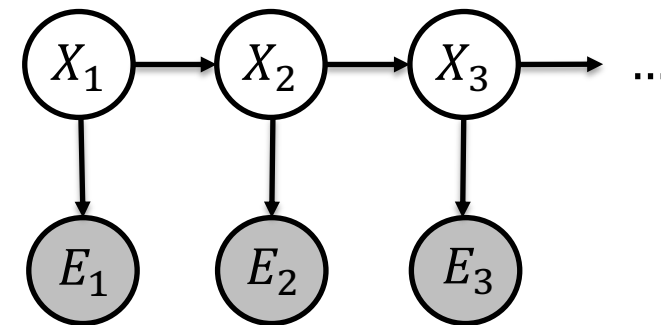
- A different task is to provide the most likely explanation

<https://tutores.com>

WeChat: cstutorcs

- Considering all possible state combinations, which one has the highest probability considering the evidence
 - Therefore, we want to compute

$$\operatorname{argmax}_{x_{1:t}} P(x_{1:t} | e_{1:t})$$



State Trellis

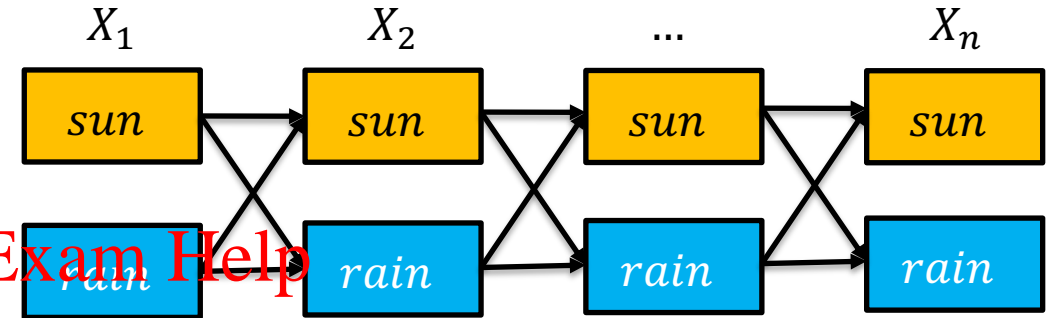
- A state trellis is a graph that illustrates the state transition over time

- Each arc represents a time passage/evidence observation with weight

$$P(x_t|x_{t-1})P(e_t|x_t)$$

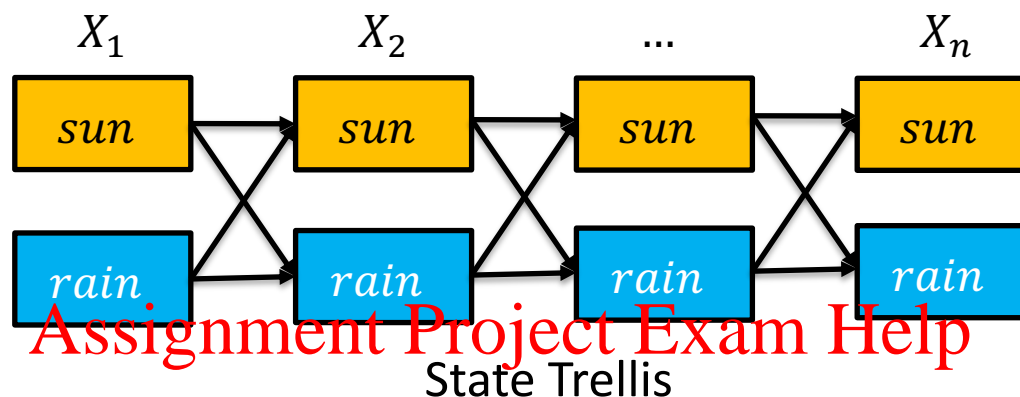
- A *path* is a sequence of states

- The product of weights on a path is the sequence probability according to the evidence
- The forward algorithm computes sums of paths probabilities that end in a same state, such as $X_n = \text{sun}$
- We will see now the *Viterbi algorithm* that computes the path with highest probability



State Trellis

Forward and Viterbi Algorithms



<https://tutorcs.com>

WeChat: cstutorcs

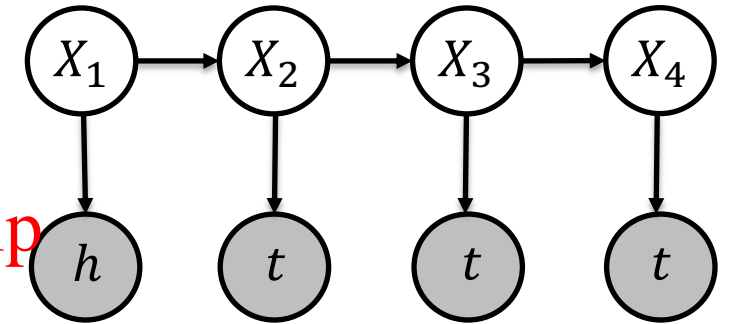
- The forward algorithm computes the *sum* of the path probabilities that lead to the same final state
- The Viterbi algorithm computes the *maximum* of the path probabilities that lead to the same final state

$$\begin{aligned}
 s[x_t] &= P(x_t | e_{1:t}) \\
 &= P(e_t | x_t) \sum_{x_{t-1}} P(x_t | x_{t-1}) s[x_{t-1}]
 \end{aligned}$$

$$\begin{aligned}
 m[x_t] &= \max_{x_{1:t-1}} P(x_{1:t-1}, x_t | e_{1:t}) \\
 &= P(e_t | X_t) \max_{x_{t-1}} P(x_t | x_{t-1}) m[x_{t-1}]
 \end{aligned}$$

Viterbi Algorithm

- Consider we have two unfair coins, c_1 and c_2
 - Someone flips the coins sequentially, but we do not know which one. We only observe the outcomes *heads* or *tails*
 - But we know that c_1 has a higher probability of heads and c_2 of tails
 - Also, the person has a preference to keep the same coin and he/she starts with a coin chosen randomly with equal probabilities

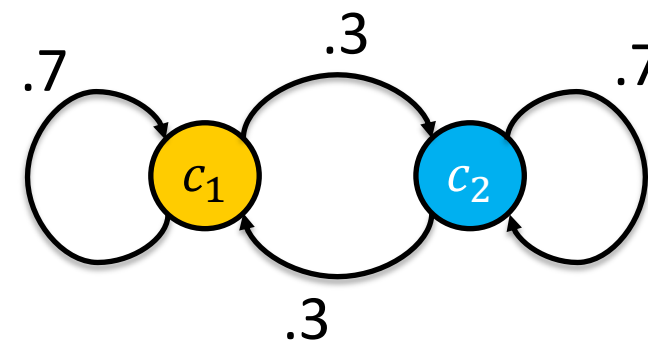
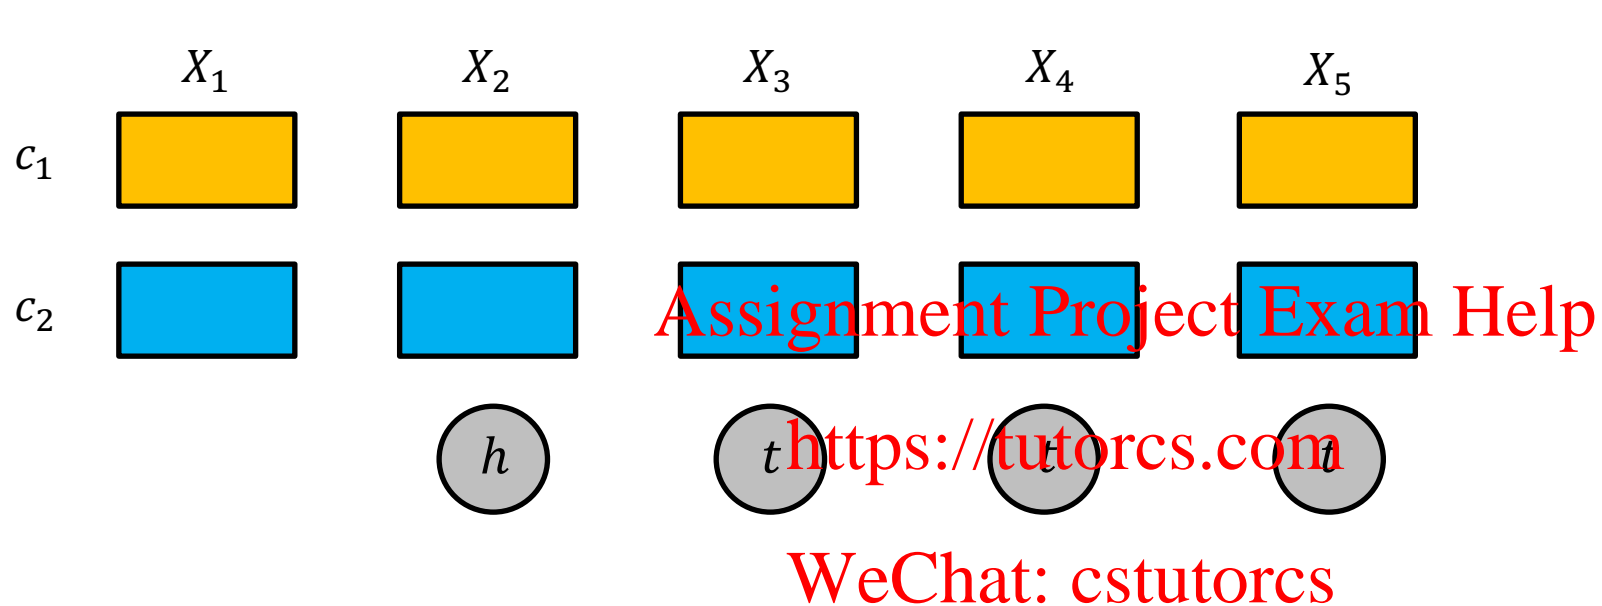


X_1	$P(X_1)$
c_1	.5
c_2	.5

X_{t-1}	X_t	$P(X_t X_{t-1})$
c_1	c_1	.7
c_1	c_2	.3
c_2	c_1	.3
c_2	c_2	.7

X_t	E_t	$P(E_t X_t)$
c_1	h	.8
c_1	t	.2
c_2	h	.2
c_2	t	.8

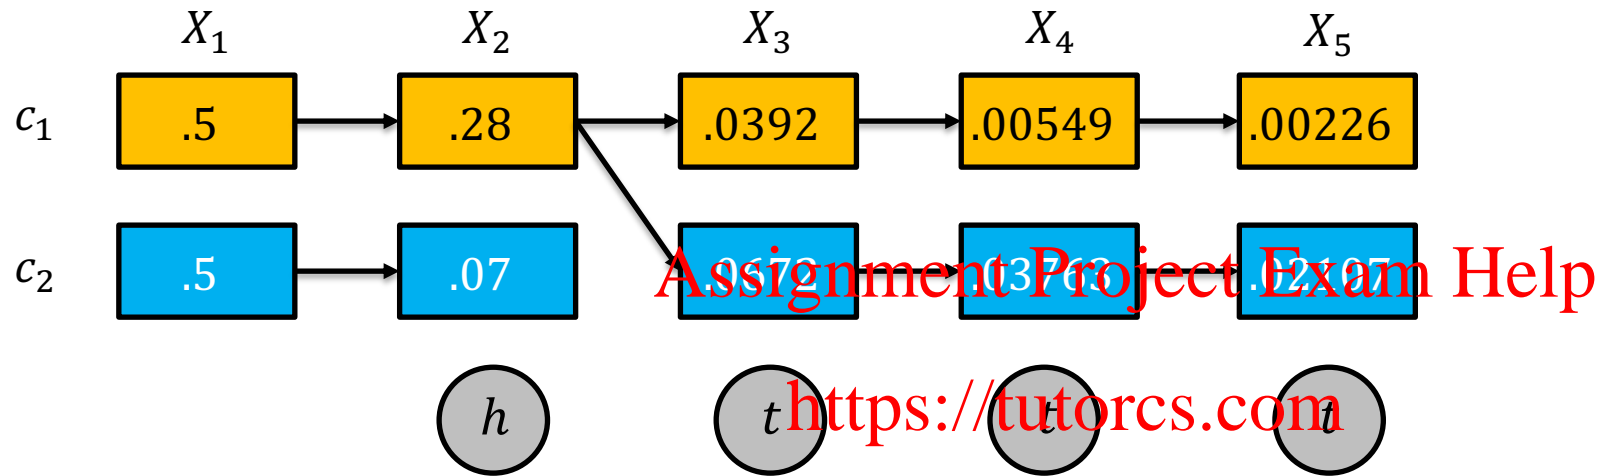
Viterbi Algorithm



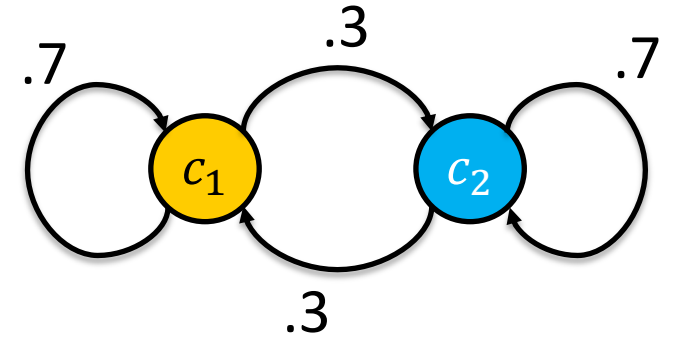
X_1	$P(X_1)$	X_t	E_t	$P(E_t X_t)$
c_1	.5	c_1	h	.8
c_2	.5	c_1	t	.2
		c_2	h	.2
		c_2	t	.8

$$\begin{aligned}
 m[x_t] &= \max_{x_{1:t-1}} P(x_{1:t-1}, x_t | e_{1:t}) \\
 &= P(e_t | X_t) \max_{x_{t-1}} P(x_t | x_{t-1}) m[x_{t-1}]
 \end{aligned}$$

Viterbi Algorithm



WeChat: cstutorcs



X_1	$P(X_1)$	X_t	E_t	$P(E_t X_t)$
c_1	.5	c_1	h	.8
c_2	.5	c_1	t	.2
		c_2	h	.2
		c_2	t	.8

$$\begin{aligned}
 m[x_t] &= \max_{x_{1:t-1}} P(x_{1:t-1}, x_t | e_{1:t}) \\
 &= P(e_t | X_t) \max_{x_{t-1}} P(x_t | x_{t-1}) m[x_{t-1}]
 \end{aligned}$$

Viterbi Algorithm

Input: time n , transition probability T , emission probability E , prior probability of states $P(X_1)$, sequence of observations $\{e_2, \dots, e_t\}$

Output: $\max_{x_{1:t-1}} P(x_{1:t-1}, x_t | e_{2:t})$

for each state x **do**

$m[x, 1] \leftarrow P(X_1 = x)$

for $t \leftarrow 2$ to n **do**

for each state x_t **do**

$m[x_t, t] = 0$

for each state x_{t-1} **do**

if $m[x_{t-1}, t-1]T(x_t|x_{t-1}) > m[x_t, t]$

$m[x_t, t] \leftarrow m[x_{t-1}, t-1]T(x_t|x_{t-1})$

$m[x_t, t] \leftarrow m[x_t, t]E(e_t|x_t)$

return $p[x, n]$ for all states x

$$O(n|X|^2)$$

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

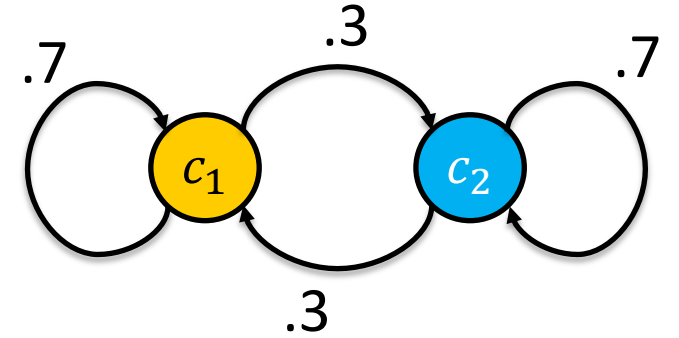
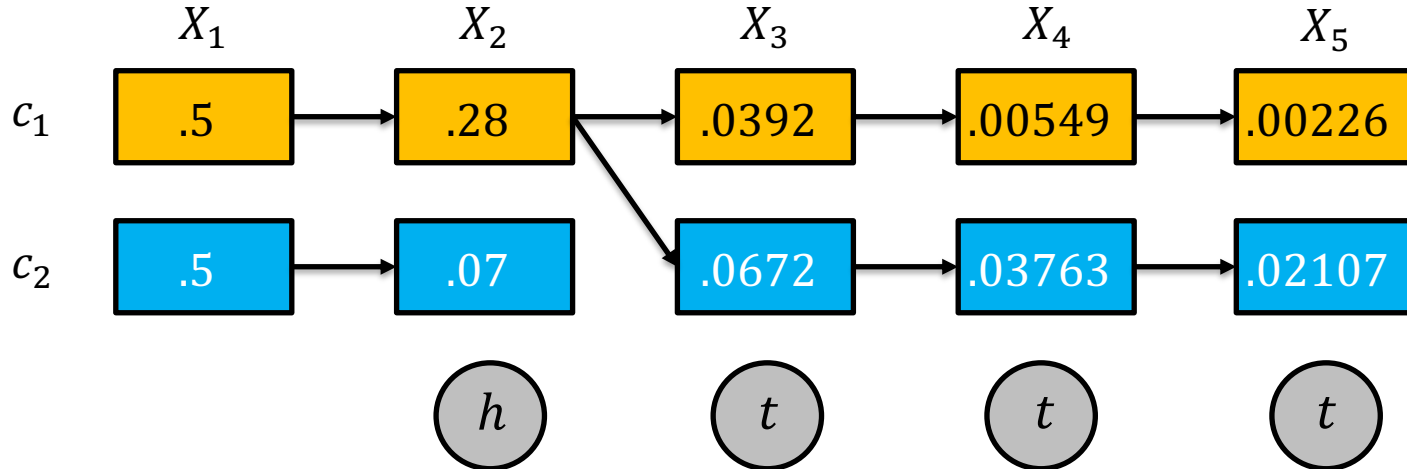
Viterbi Algorithm

- The Viterbi algorithm of the previous slide provides the probability of the most likely sequence
 - However, often we are more interested in the sequence instead of its probability
- There are two common solutions
 - Keep an additional structure pointing to the parent of each node
 - Backtrack the computation from the last node

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: estutorcs

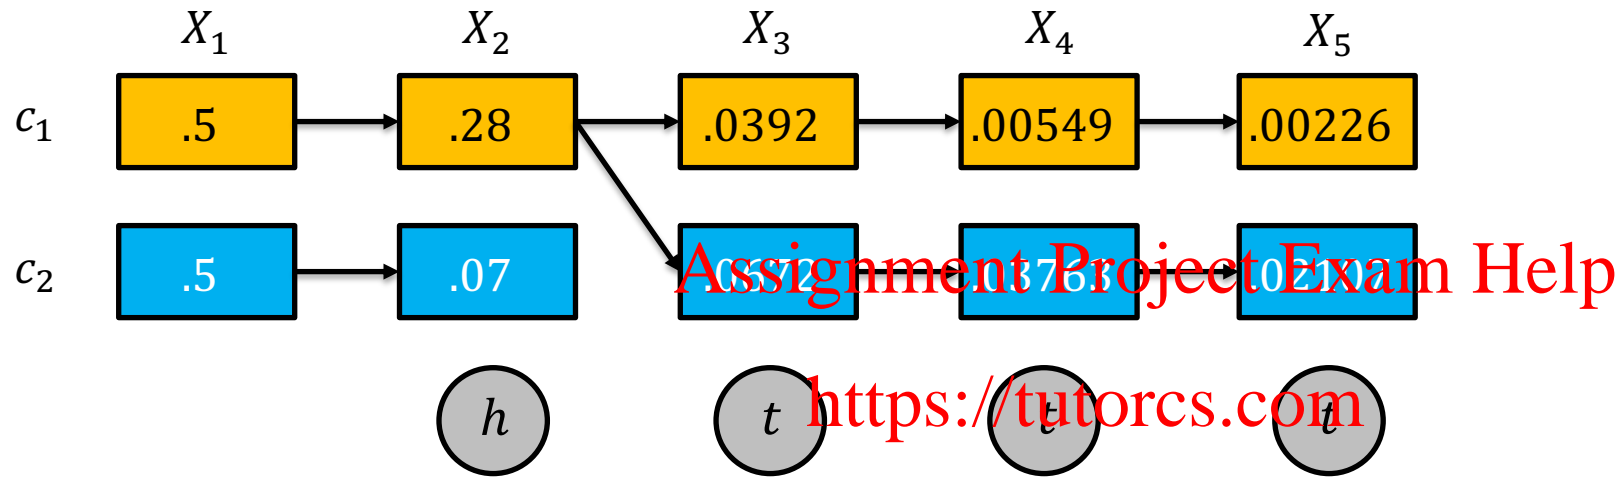


X_1	$P(X_1)$	X_t	E_t	$P(E_t X_t)$
c_1	.5	c_1	h	.8
c_2	.5	c_1	t	.2
		c_2	h	.2
		c_2	t	.8

$$m[x_t] = \max_{x_{1:t-1}} P(x_{1:t-1}, x_t | e_{1:t})$$

$$= P(e_t | X_t) \max_{x_{t-1}} P(x_t | x_{t-1}) m[x_{t-1}]$$

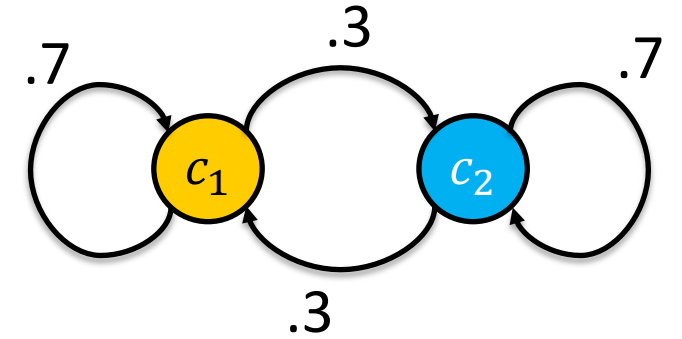
Viterbi Algorithm: Backtracking Computation



Repeat

1. Divide by the probability of evidence
2. For each state x_{t-1} divide by $P(x_t|x_{t-1})$
3. See which value of x_{t-1} matches the result

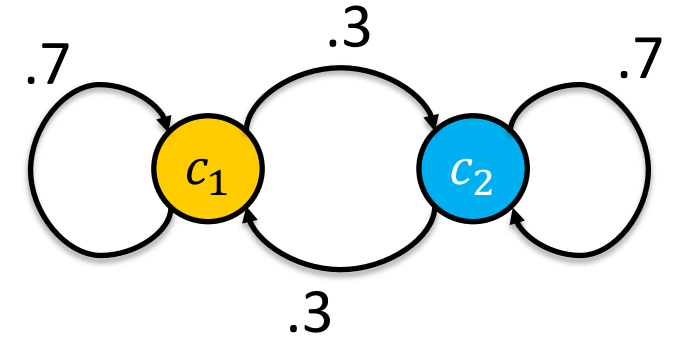
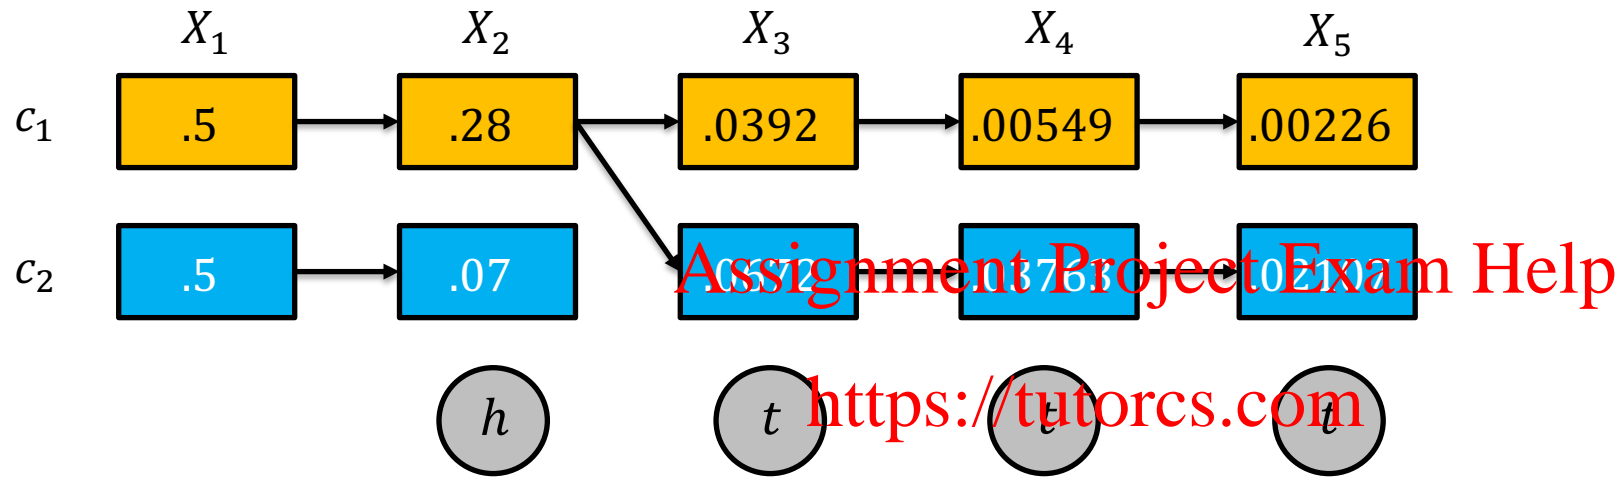
WeChat: cstutorcs



X_1	$P(X_1)$	X_t	E_t	$P(E_t X_t)$
c_1	.5	c_1	h	.8
c_2	.5	c_1	t	.2
		c_2	h	.2
		c_2	t	.8

$$\begin{aligned}
 m[x_t] &= \max_{x_{1:t-1}} P(x_{1:t-1}, x_t | e_{1:t}) \\
 &= P(e_t | X_t) \max_{x_{t-1}} P(x_t | x_{t-1}) m[x_{t-1}]
 \end{aligned}$$

Viterbi Algorithm: Backtracking Computation



X_1	$P(X_1)$	X_t	E_t	$P(E_t X_t)$
c_1	.5	c_1	h	.8
c_2	.5	c_1	t	.2
		c_2	h	.2
		c_2	t	.8

Repeat

1. Divide by the probability of evidence
2. For each state x_{t-1} divide by $P(x_t|x_{t-1})$
3. See which value of x_{t-1} matches the result

1.

$$\frac{.02107}{.8} = 0.0263375$$

2.

$$x_4 = c_1: \frac{0.0263375}{.3} \approx 0.08779$$

$$x_4 = c_2: \frac{0.0263375}{.7} \approx 0.03763$$

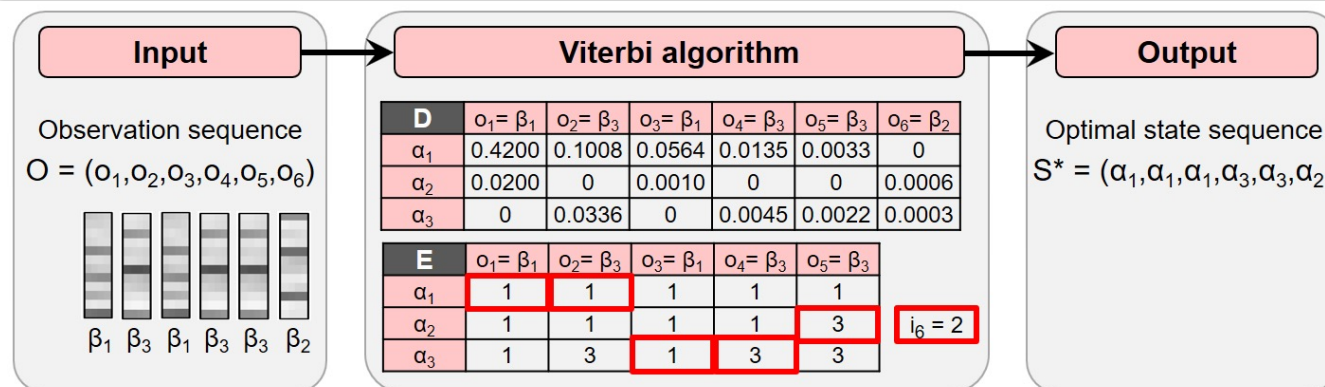
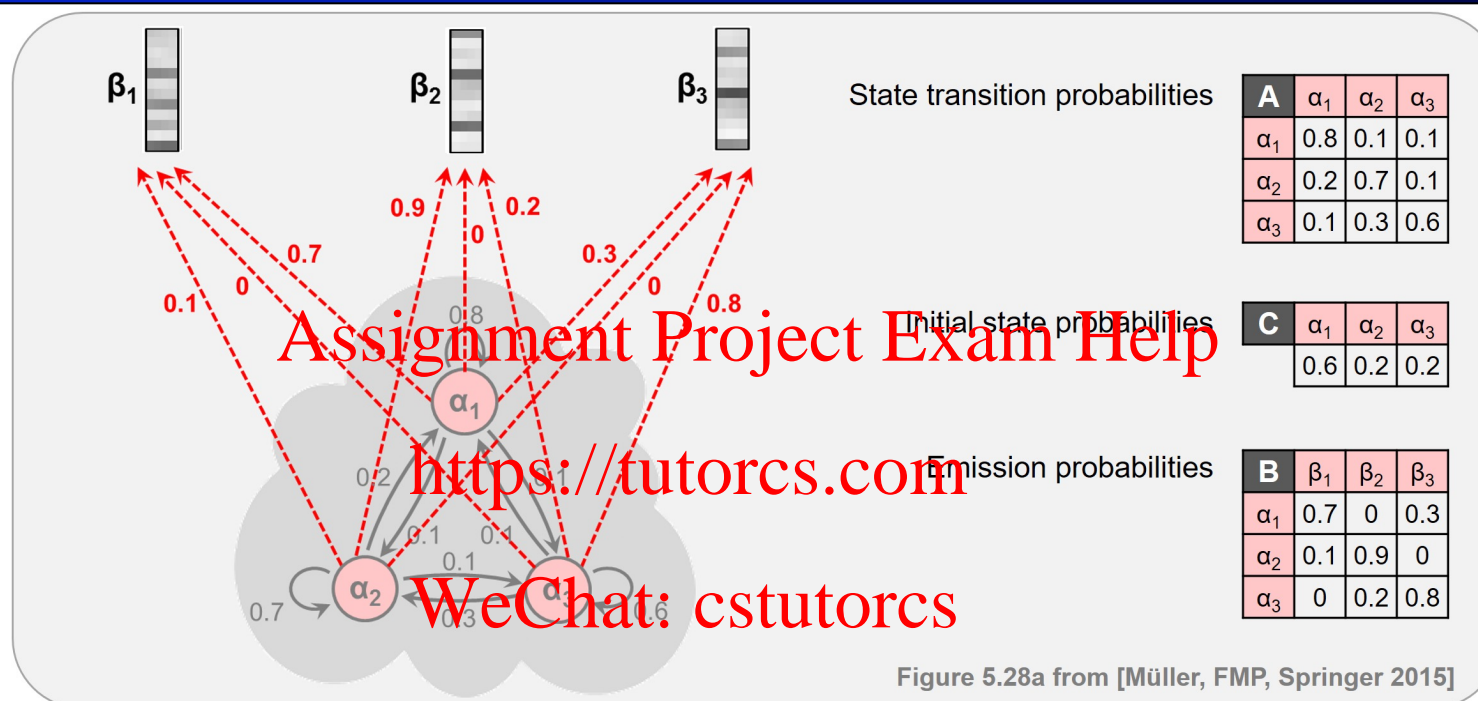
3.

$$x_4 = c_2: 0.03763$$

$$m[x_t] = \max_{x_{1:t-1}} P(x_{1:t-1}, x_t | e_{1:t})$$

$$= P(e_t | X_t) \max_{x_{t-1}} P(x_t | x_{t-1}) m[x_{t-1}]$$

Viterbi Algorithm: Backtracking Computation



Viterbi Algorithm: Vanishing Probabilities

- Notice the probabilities decrease as we observe more evidence
 - It is intuitive since the number of paths grows exponentially with the sequence size
 - In this example, the probabilities are around 10^{-4} with just 6 steps
 - Long sequences (such as 100 steps) will cause an underflow and the probabilities will become zero
 - We can fix that using log probabilities, similarly to the Naïve Bayes classifier
 - This approach also replaces multiplications by sums

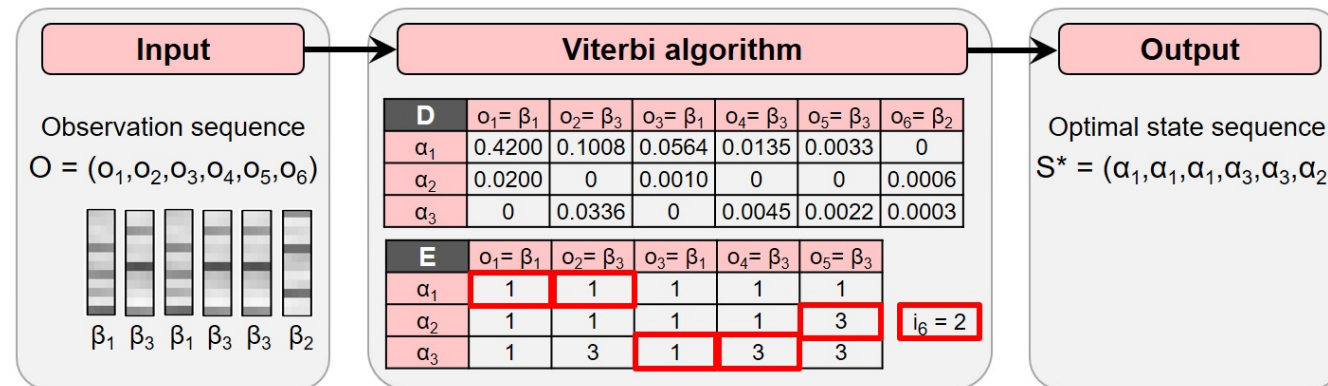


Figure 5.28b from [Müller, FMP, Springer 2015]

Viterbi Algorithm: Log Probabilities

Input: time n , transition probability T , emission probability E , prior probability of states $P(X_1)$, sequence of observations $\{e_2, \dots, e_t\}$

Output: $\max_{x_{1:t-1}} \log P(x_{1:t-1}, x_t | e_{1:t})$

for each state x do

$m[x, 1] \leftarrow \log P(X_1 = x)$

for $t \leftarrow 2$ to n do

for each state x_t do

$m[x_t, t] = -\infty$

for each state x_{t-1} do

if $m[x_{t-1}, t-1] + \log T(x_t | x_{t-1}) > m[x_t, t]$

$m[x_t, t] \leftarrow m[x_{t-1}, t-1] + \log T(x_t | x_{t-1})$

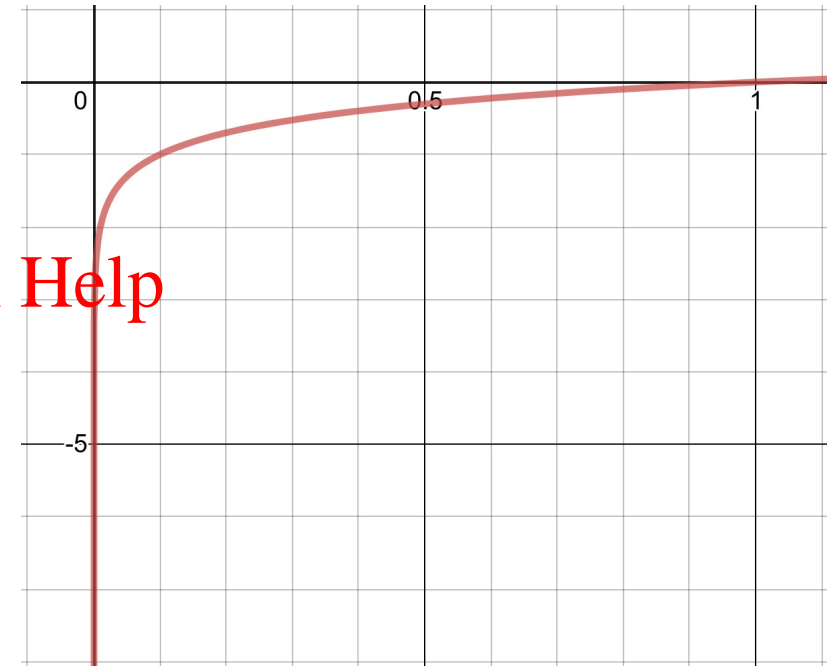
$m[x_t, t] \leftarrow m[x_t, t] + \log E(e_t | x_t)$

return $p[x, n]$ for all states x

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs



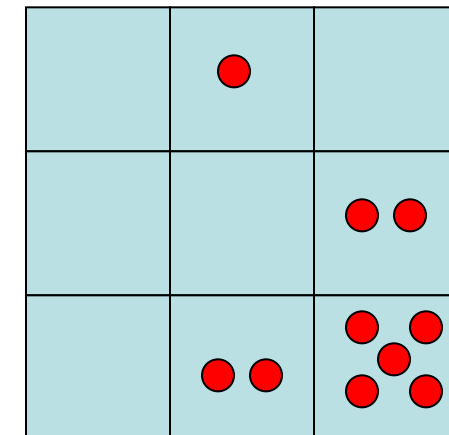
$$\begin{aligned} m[x_t] &= \log \max_{x_{1:t-1}} P(x_{1:t-1}, x_t | e_{1:t-1}) \\ &= \log P(e_t | X_t) + \max_{x_{t-1}} \log P(x_t | x_{t-1}) + m[x_{t-1}] \end{aligned}$$

$$O(n|X|^2)$$

Particle Filtering

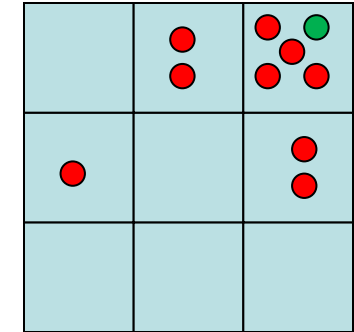
- Filtering: approximate solution
- Sometimes $|X|$ is too big to use exact inference
 - $|X|$ may be too big to even store $B(X)$
 - E.g. X is continuous
- Solution: approximate inference
 - Track samples of X , not all values
 - Samples are called particles
 - Time per step is linear in the number of samples
 - But: number needed may be large
 - In memory: list of particles, not states
- This is how robot localization works in practice
- Particle is just new name for sample

0.0	0.1	0.0
0.0	0.0	0.2
0.0	0.2	0.5



Representation: Particles

- Our representation of $P(X)$ is now a list of N particles (samples)
 - Generally, $N \ll |X|$
 - Storing map from X to counts would defeat the point
- $P(x)$ approximated by number of particles with value x
 - So, many x may have $P(x) = 0$
 - More particles, more accuracy
- For now, all particles have a weight of 1



Particles:

(1,3)
(1,2)
(1,3)
(2,3)
(1,3)
(2,3)
(2,1)
(1,3)
(1,3)
(1,2)

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

Particle Filtering: Elapse Time

- Each particle is moved by sampling its next position from the transition model

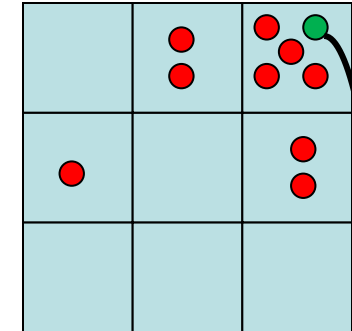
$$x' = \text{sample}(P(x'|g))$$

- Here, most samples move clockwise, but some move in another direction or stay in place

- This captures the passage of time
 - If enough samples, close to exact values before and after (consistent)

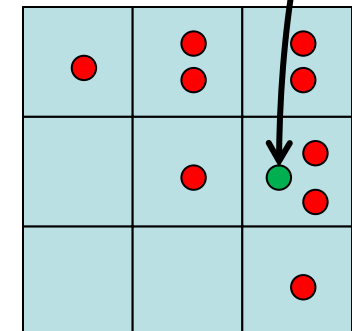
Particles:

(1,3)
(1,2)
(1,3)
(2,3)
(2,3)
(2,3)
(2,1)
(1,3)
(1,3)
(1,2)



Particles:

(2,3)
(1,2)
(2,3)
(3,3)
(1,3)
(2,3)
(1,1)
(1,2)
(1,3)
(2,2)



Particle Filtering: Observe

- Slightly trickier:

- Don't sample observation, fix it
- Downweigh samples based on the evidence

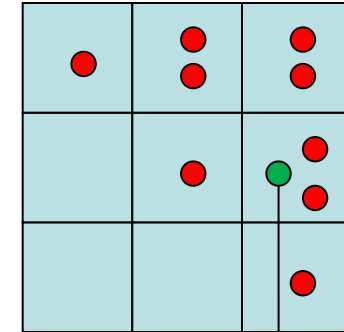
$$w(x) = P(e|x)$$

$$B(X) \propto P(e|X)B'(X)$$

- As before, the probabilities don't sum to one, since all have been downweighed (in fact they now sum to $(N \text{ times})$ an approximation of $P(e)$)

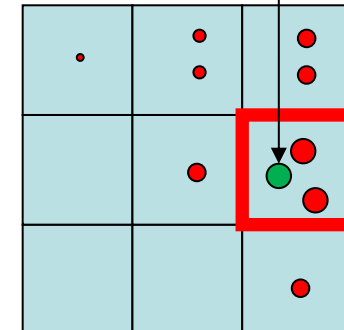
Particles:

(2,3)
(1,2)
(2,3)
(3,3)
(1,3)
(2,3)
(1,1)
(1,2)
(1,3)
(2,2)



Particles:

(2,3) w=.9
(1,2) w=.2
(2,3) w=.9
(3,3) w=.4
(1,3) w=.4
(2,3) w=.9
(1,1) w=.1
(1,2) w=.2
(1,3) w=.4
(2,2) w=.4



Particle Filtering: Resample

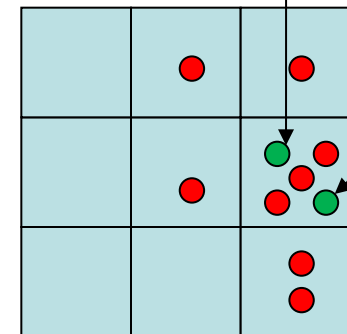
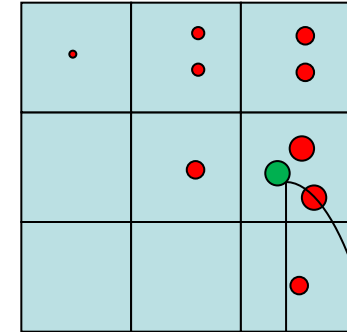
- Rather than tracking weighted samples, we resample
- N times, we choose from our weighted sample distribution (i.e. draw with replacement)
- This is equivalent to renormalizing the distribution
- Now the update is complete for this time step, continue with the next one

Particles:

(2,3) $w=.9$
(1,2) $w=.2$
(2,3) $w=.9$
(3,3) $w=.4$
(1,3) $w=.4$
(2,3) $w=.9$
(1,1) $w=.1$
(1,2) $w=.2$
(1,3) $w=.4$
(2,2) $w=.4$

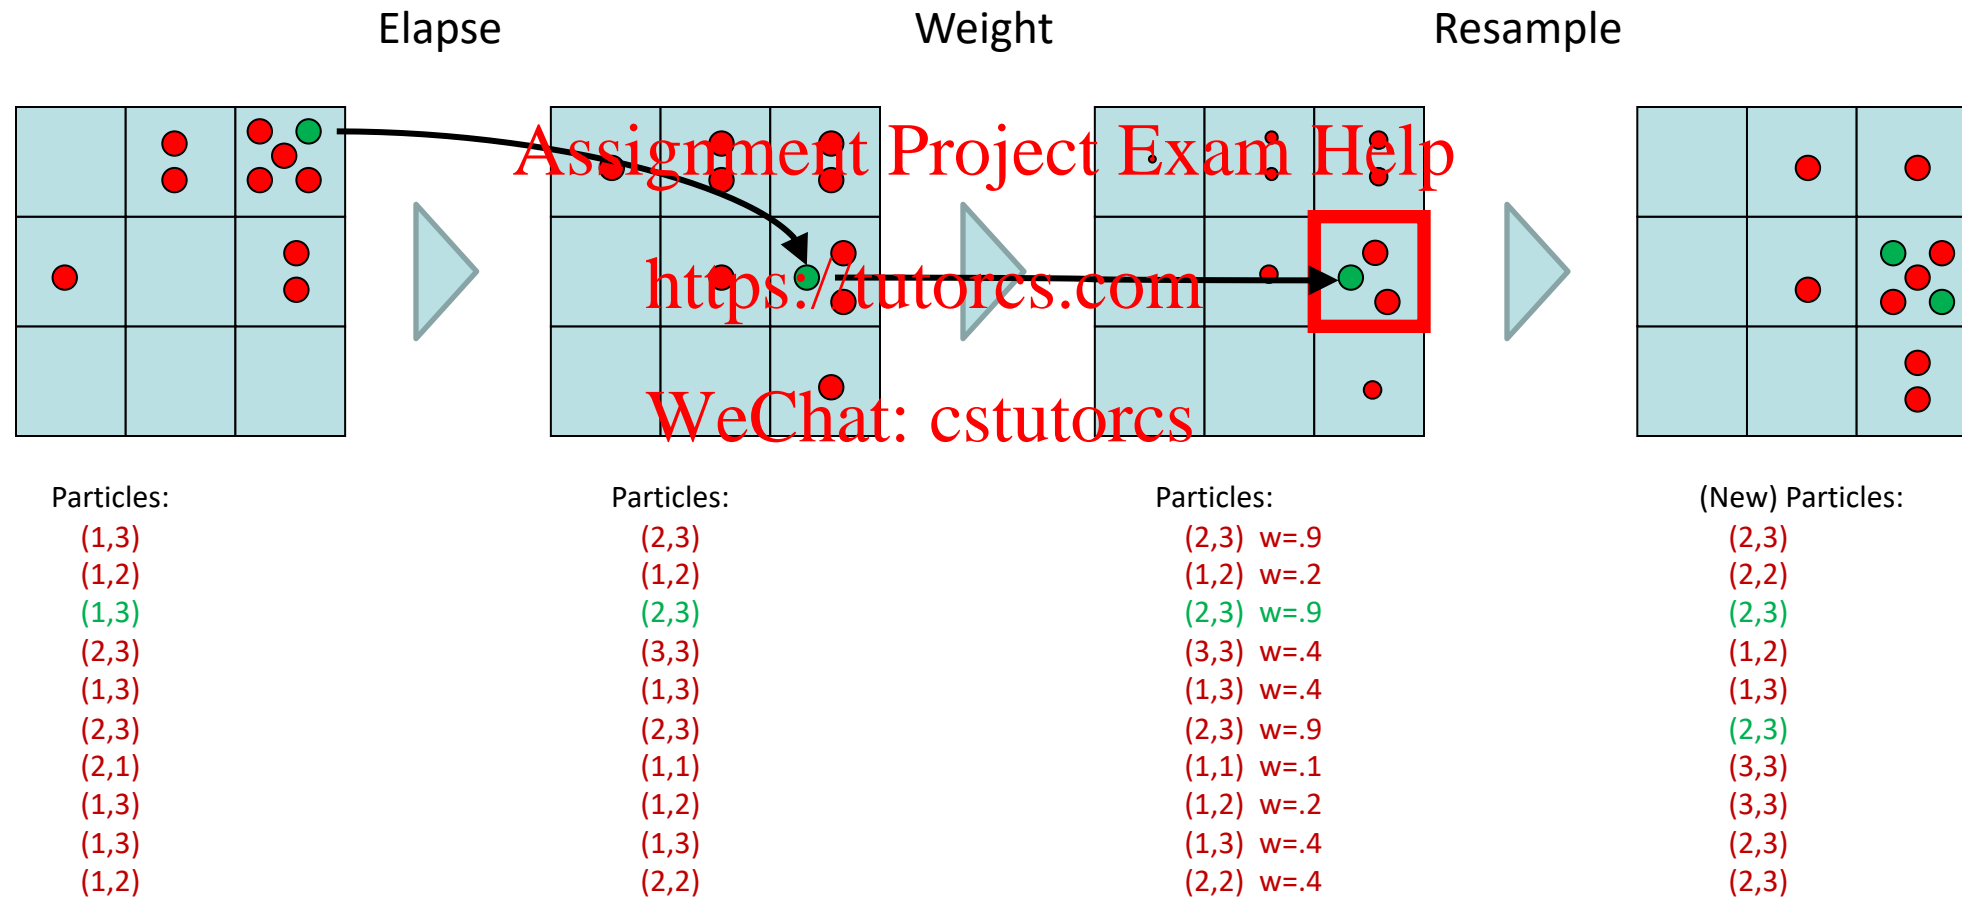
(New) Particles:

(2,3)
(2,2)
(2,3)
(1,2)
(1,3)
(2,3)
(3,3)
(3,3)
(2,3)
(2,3)



Recap: Particle Filtering

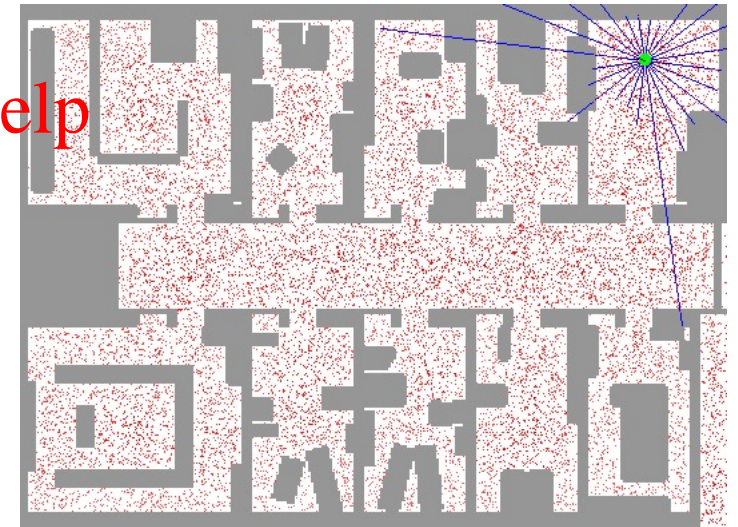
- Particles: track samples of states rather than an explicit distribution



Robot Localization

- In robot localization:

- We know the map, but not the robot's position
- Observations may be vectors of range finder readings
- State space and readings are typically continuous (works basically like a very fine grid) and so we cannot store $B(X)$
- Particle filtering is a main technique



Conclusion

- Markov chains and Hidden Markov models are simple examples of Dynamic Bayesian networks
 - DBNs are networks that allow us to model changes in time or space
 - Changes in time are specified using transition probabilities
- Markov chains are sequence models
 - It tracks the probability distribution over a series of transitions
 - For many sequences, the probability distribution converges to a stationary distribution
 - The stationary distribution has several applications such as the MCMC algorithms used for approximate inference
- Hidden Markov models are Markov chains with hidden states
 - Those states are never directly observed, but we can make indirect inference through emissions
 - These models are used in several applications such as language and signal processing and robot localisation